


NOW AND THEN, I ANNOUNCE "I KNOW
YOU'RE LISTENING" TO EMPTY ROOMS.



IF I'M WRONG, NO ONE KNOWS.
AND IF I'M RIGHT, MAYBE I JUST FREAKED
THE HELL OUT OF SOME SECRET ORGANIZATION.


CS 60: Computer Networks

Packet sniffing and spoofing

Review from last class

1. All hosts on the Internet have an IP address
2. How does a host get an IP address?
 - Statically (manually) assigned by network administrator
 - DHCP server provides IP address from a pool of addresses
3. Network Interface Cards (NICs) have a unique Media Access Controller (MAC) address burned into ROM
4. Switches direct traffic on a LAN based on MAC address
5. How does a computer find the MAC of another computer?
 - ARP
6. What's left?
 - Routing between LANs
7. Before we get to routing, we'll take a closer look at higher layers then we will focus on creating and sniffing network traffic

Agenda

- 
1. Review: network layers
 2. Sockets
 3. Scapy
 4. Exercises

Review: network layers

Network model

Critical information

7) Application

Application data to be sent across network

Data

4) Transport

3) Network (IP)

2) Link (MAC)

1) Physical

Review: network layers

Network model

Critical information

7) Application

Application data to be sent across network

4) Transport

TCP or UDP


**TCP is connection oriented,
missing packets are resent**

3) Network (IP)

**UDP is not connection oriented,
missing packets are not resent**

2) Link (MAC)

1) Physical



Demo: UDP vs TCP



Write message on pieces of paper

Send paper to a student volunteer

- UDP may drop packets, and they may also arrive out of order
- TCP will ACK packets to make sure missing packets are resent, and packets can be put in proper order

At the Transport Layer, TCP is connection-oriented, UDP is not


Transmission Control Protocol (TCP)



We will look at the transport layer in detail soon

User Datagram Protocol (UDP)

UDP does not establish a connection



UDP Transmission Illustration

- Messages are broken into packets
- Each packet given sequence number
- Packets reassembled in seq order
- Missing packets are resent

- Messages are broken into packets
- Packet received or not, no resend
- Faster than TCP, smaller headers
- Use for streaming

Layer 4 ports identify an application running on a host

Ports

Your computer is running multiple applications at the same time (email, web browsing, others)

Without a means to differentiate, all received data would go to the same place (e.g., web browser gets email traffic and vice versa)

Ports identify applications by a number that ranges from 0 to $65,535 = 2^{16}-1$

Servers run applications that listen for connections on well-known ports

Clients connect to servers on these ports

Well known ports	
Port	Service
20 and 21	FTP
22	SSH
23	Telnet
53	DNS
80	HTTP
443	Encrypted HTTP
587 (old 25)	Email (SMTP)

Ports

0 - 1023 are well-known ports

Commonly reserved for system apps

1024-49,151 user or registered ports

49,152-65,535 ephemeral ports

See list at </etc/services>

Review: network layers

Network model

7) Application

4) Transport

3) Network (IP)

2) Link (MAC)

1) Physical

Critical information

Application message to be sent across network

Data

Payload is
Layer 7 data

Review: network layers

Network model

7) Application

Application message to be sent across network

4) Transport

TCP or UDP
src and dst port

3) Network (IP)

2) Link (MAC)

1) Physical

Critical information

Layer 4 uses ports to identify applications


- **A port is just an integer 0 – 65,535**

Server side:

- **Ports identify an application**
- **Different apps listen on different ports (e.g., web on port 80, telnet on port 23)**

Client side:

- **Ports identify an instance of an application (might have multiple copies of an app running)**
- **Port numbers are randomly chosen by OS**



Review: network layers

Network model

Critical information

7) Application

Application message to be sent across network

4) Transport

TCP or UDP
src and dst port

3) Network (IP)


src and dst
IP address

Routing!

2) Link (MAC)

**Packets routed across LANs
using IP addresses**

1) Physical



Review: network layers

Network model

Critical information

7) Application

Application message to be sent across network

4) Transport

TCP or UDP
src and dst port

3) Network (IP)

src and dst
IP address
Routing!

2) Link (MAC)


src and dst MAC
address

Switching!

**Frames transmitted
within a LAN based
on MAC address**

1) Physical

Source MAC Destination MAC



Review: network layers

Network model

Critical information

7) Application

Application message to be sent across network

4) Transport

TCP or UDP
src and dst port

3) Network (IP)


src and dst
IP address
Routing!

2) Link (MAC)

src and dst MAC
address
Switching!

1) Physical

Sends frames to another host over cable or RF





Sender's data goes from Layer 7 to Layer 1

Client



Sending from Client to Server

Server



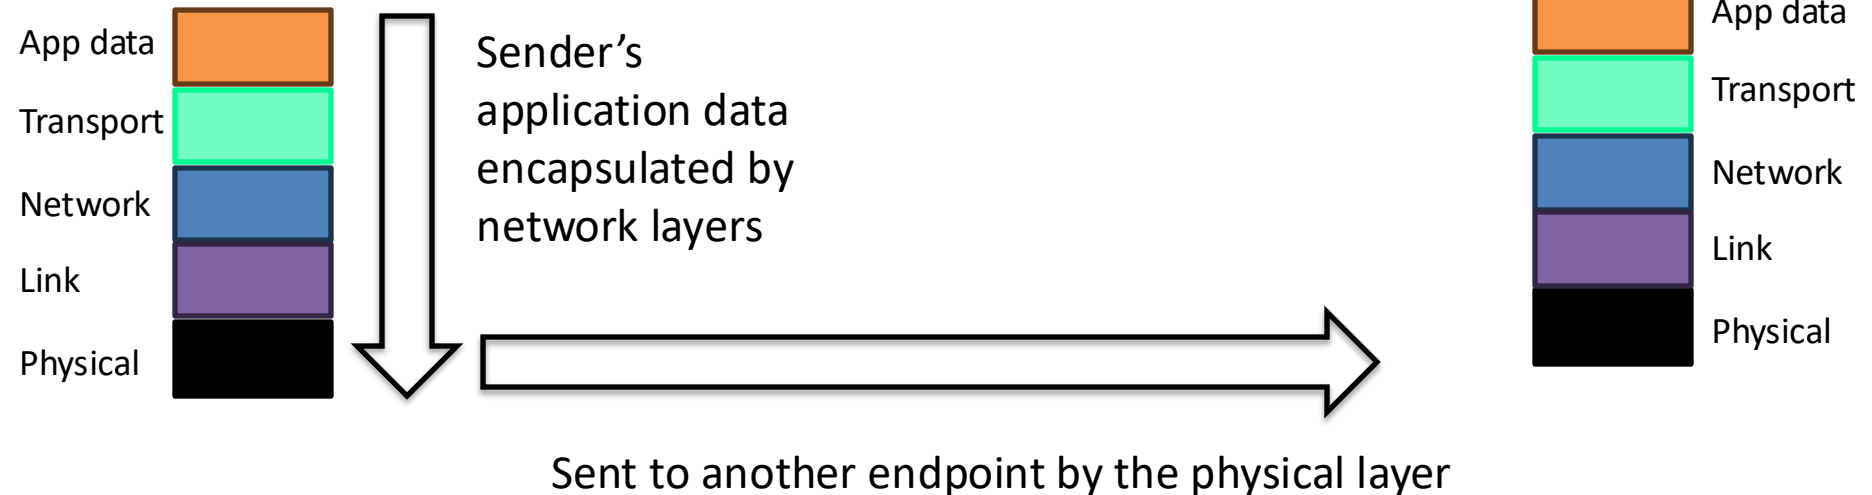
Data sent across the network on Layer 1

Client

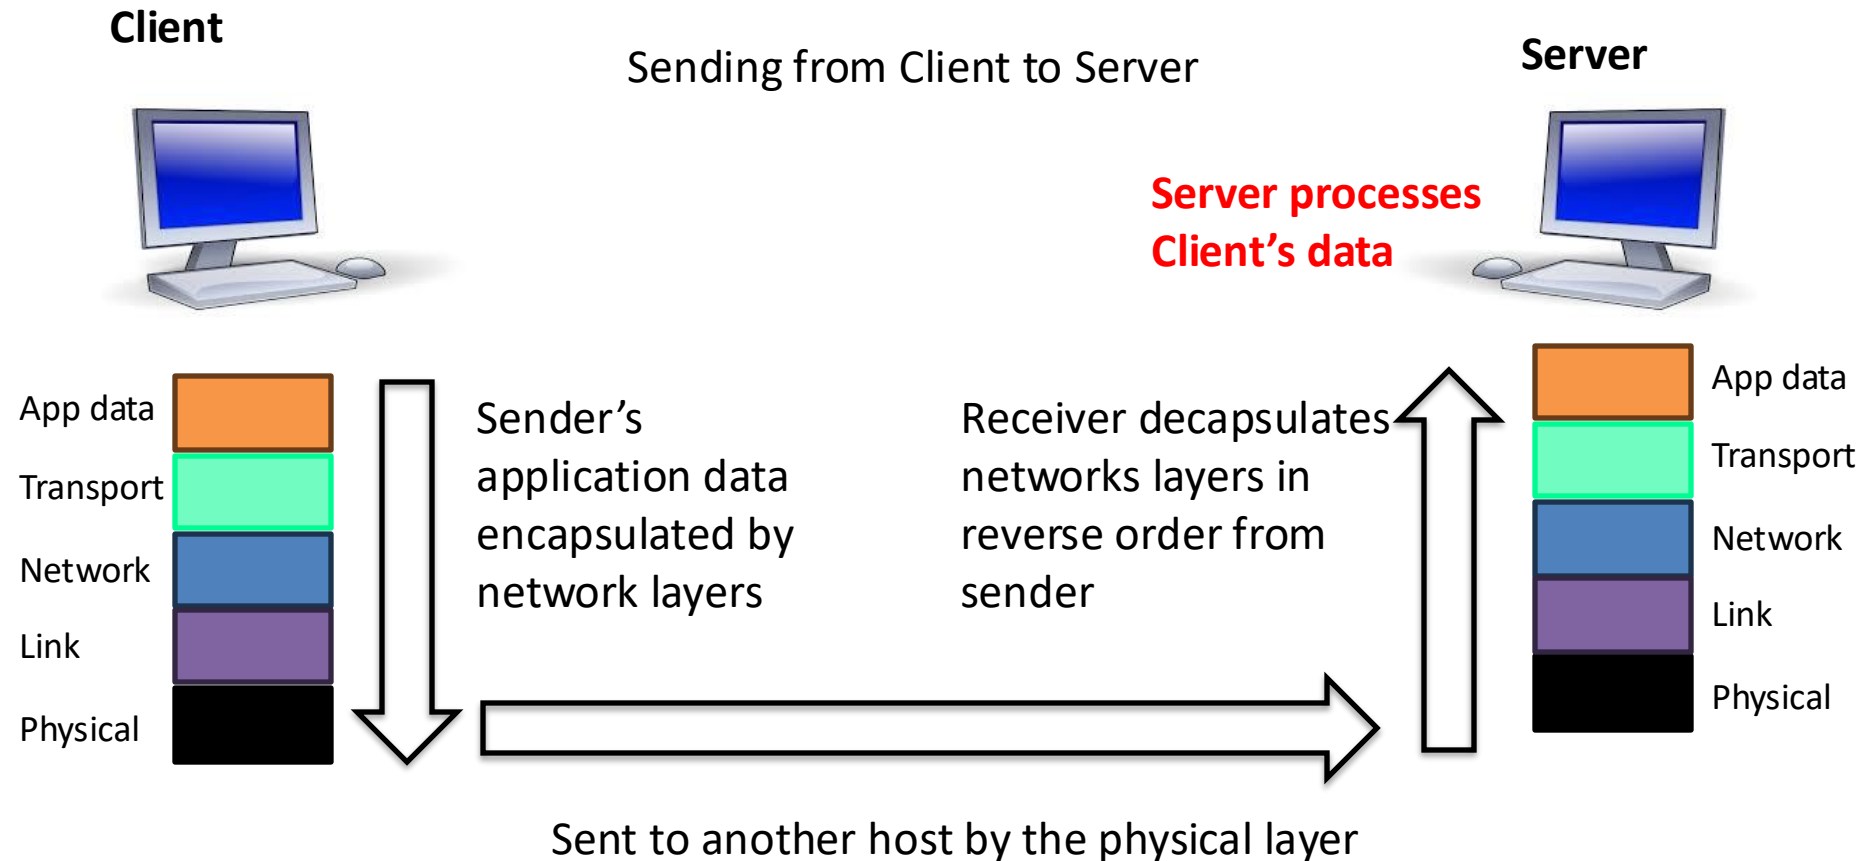


Sending from Client to Server

Server




Receiver strips off headers from Layer 2 up to Layer 7




Reply from Server reverses the process back to the Client

Client




Reply is the reverse from
Server back to Client

Server



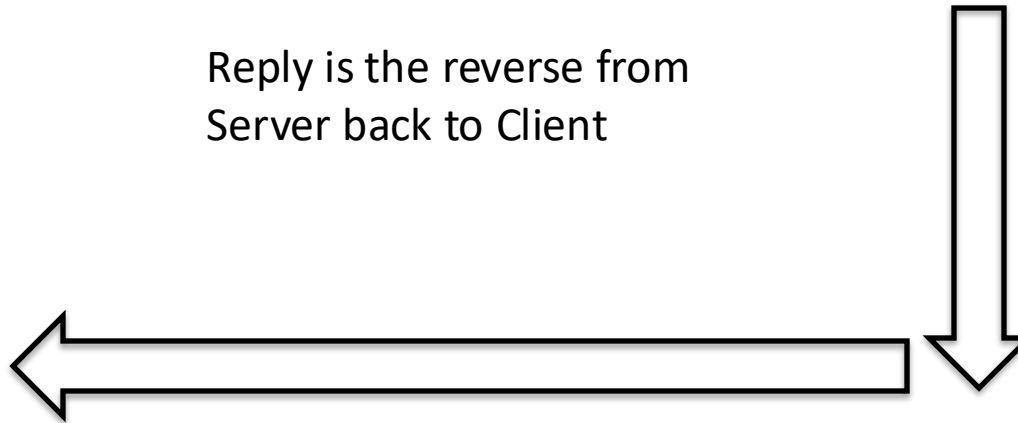

Reply from Server reverses the process back to the Client

Client



Reply is the reverse from
Server back to Client

Server



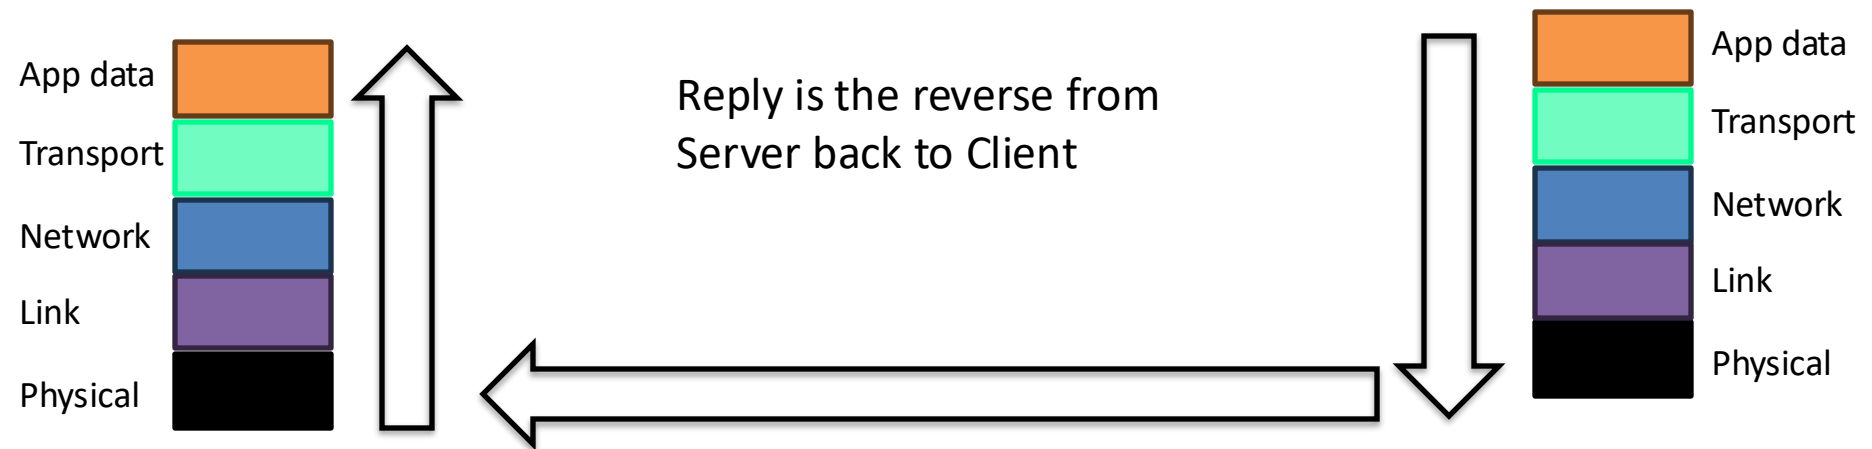
Reply from Server reverses the process back to the Client

Client



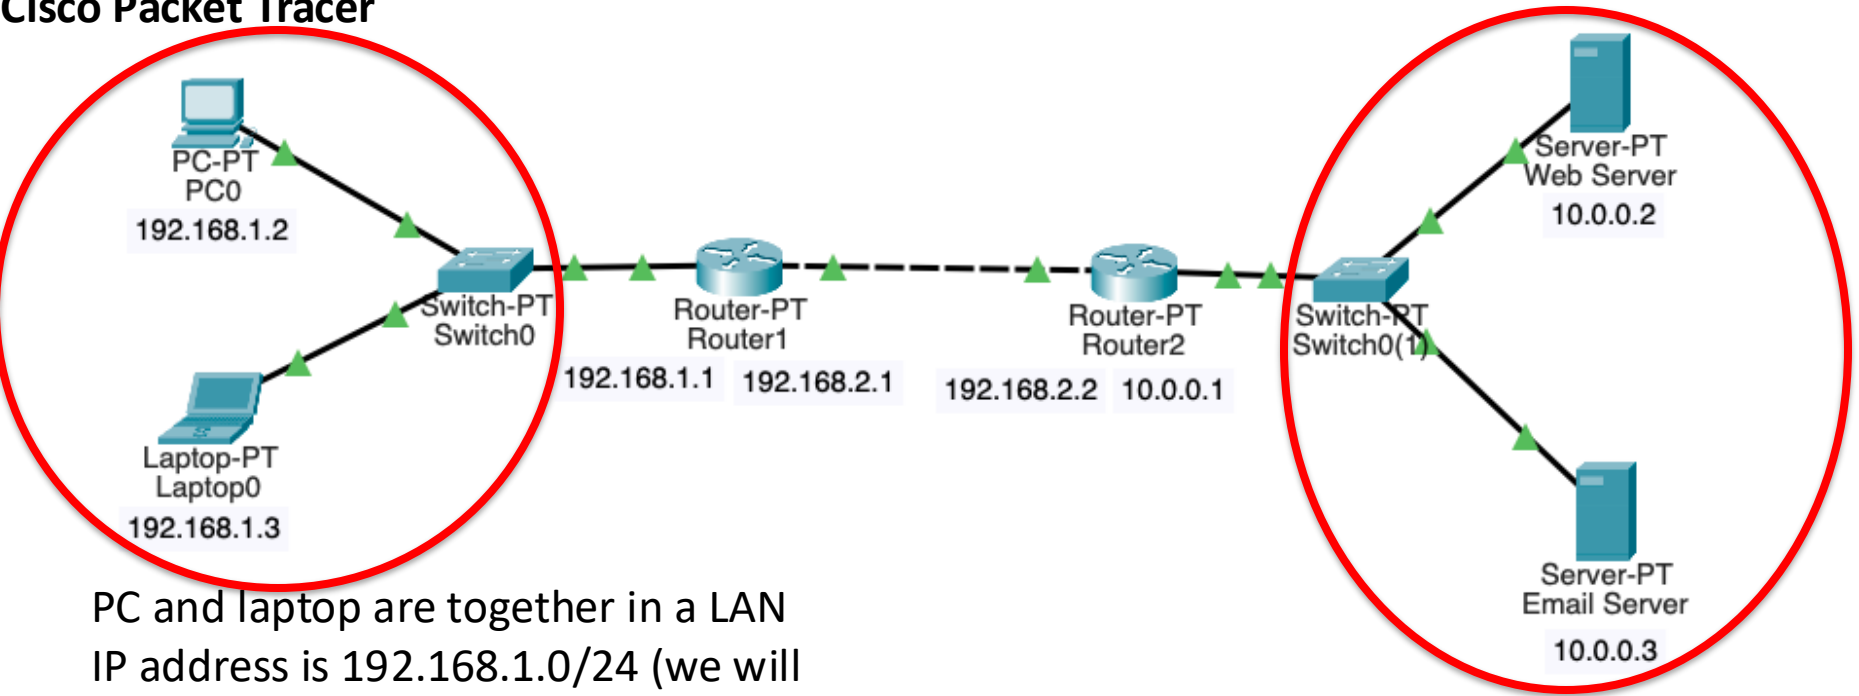
**Client processes
Server's reply**

Server



Example networks, PCs in one network, servers in another network

Cisco Packet Tracer




PC and laptop are together in a LAN
IP address is 192.168.1.0/24 (we will cover this notation soon)
This LAN is connected by a Switch

Web server and Email server are together in another LAN
IP address is 10.0.0.0/24
This LAN is connected by a different Switch

Example networks, PCs in one network, servers in another network

Cisco Packet Tracer




LANs are connected routers
Assume the Internet is in between

PC and laptop are together in a LAN
IP address is 192.168.1.0/24 (we will cover this notation soon)
This LAN is connected by a Switch

Web server and Email server are together in another LAN
IP address is 10.0.0.0/24
This LAN is connected by a different Switch

Internet largely uses a client/server model where clients make requests to servers


Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer




Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Data

Layer 7 (App): Laptop formats message to Web Server identifying the web page the Laptop wants to see


Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

sport = 1026 dport = 80




Layer 4 (Transport): Laptop sets

- dst port = 80 (well-known HTTP)
- src port = 1026 (random unused)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer




Client side chooses random unused port
If client has multiple browsers (or tabs) open, each gets a unique port, so data is sent to right app

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)



sport = 1026 dport = 80




Layer 4 (Transport): Laptop sets

- **dst port = 80 (well-known HTTP)**
- **src port = 1026 (random unused)**

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

src = 192.168.1.3 dst = 10.0.0.2

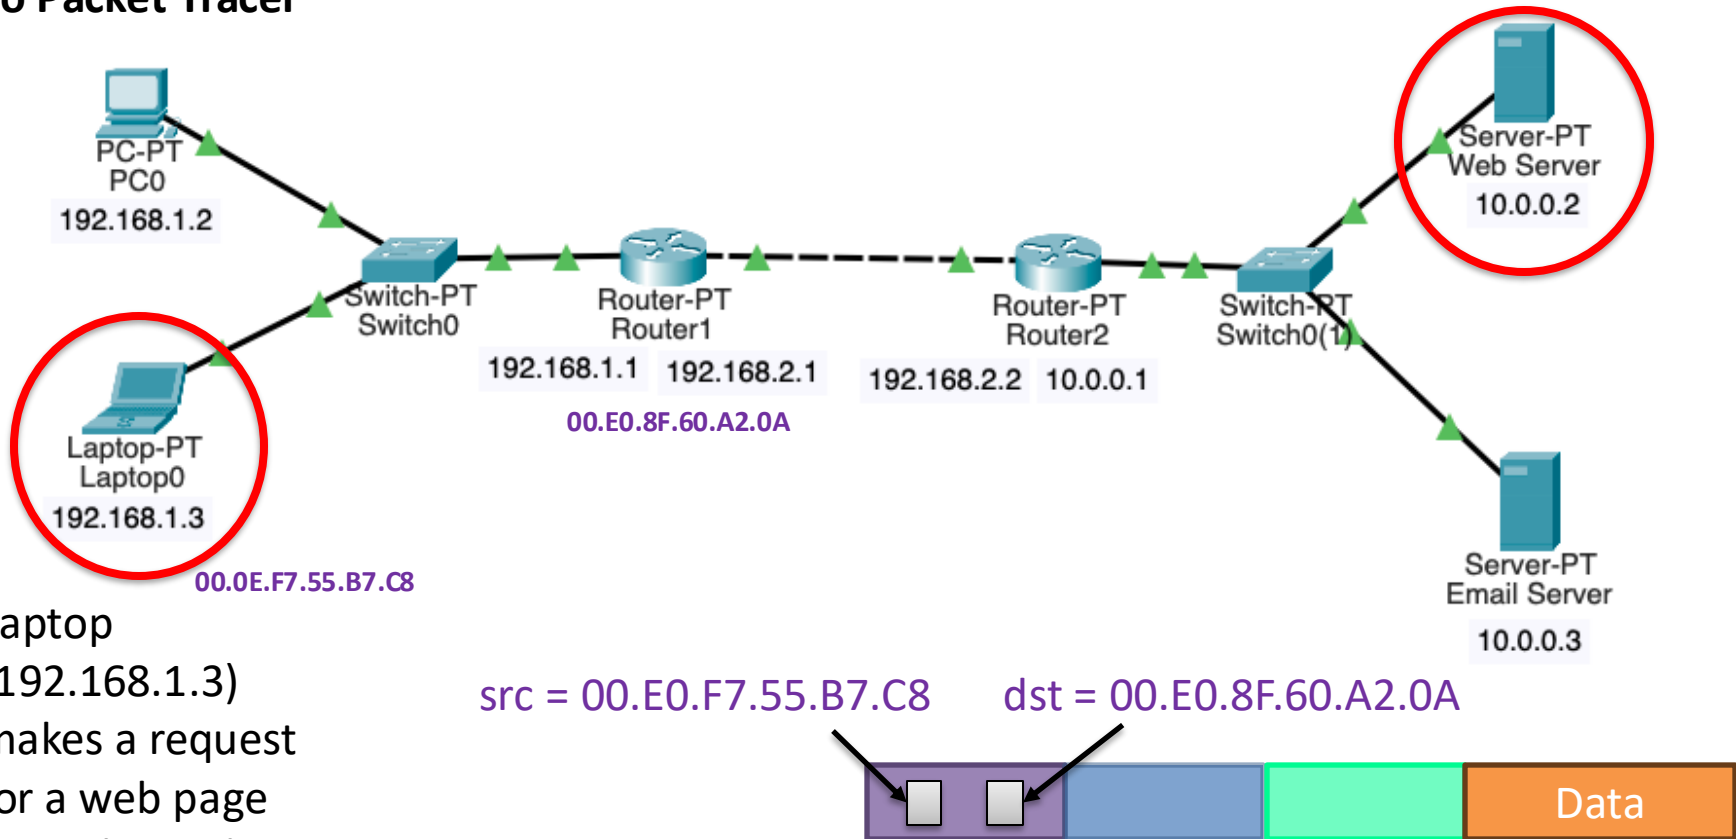


Layer 3 (Network): Laptop sets IP addresses

- **dst= 10.0.0.2 (Web server)**
- **src = 192.168.1.3 (Laptop)**

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



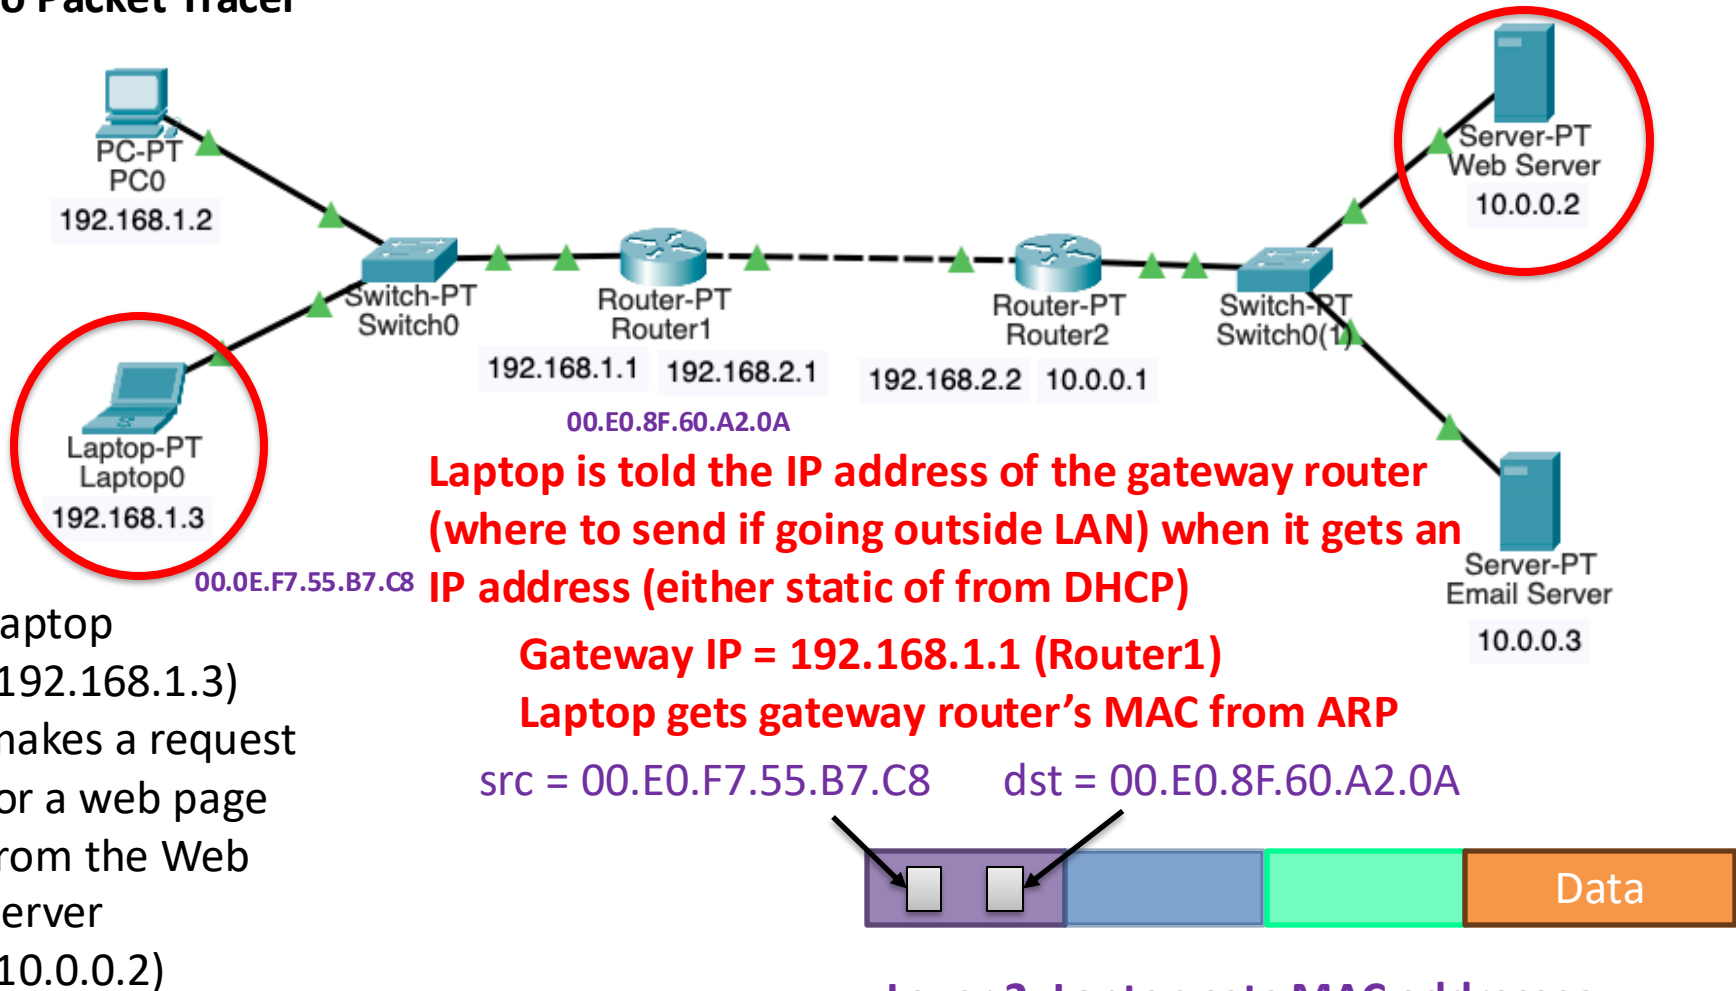
Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Layer 2: Laptop sets MAC addresses

- dst= 00.E0.8F.60.A2.0A (Router 1)
- src = 00.E0.F7.55.B7.C8 (Laptop)

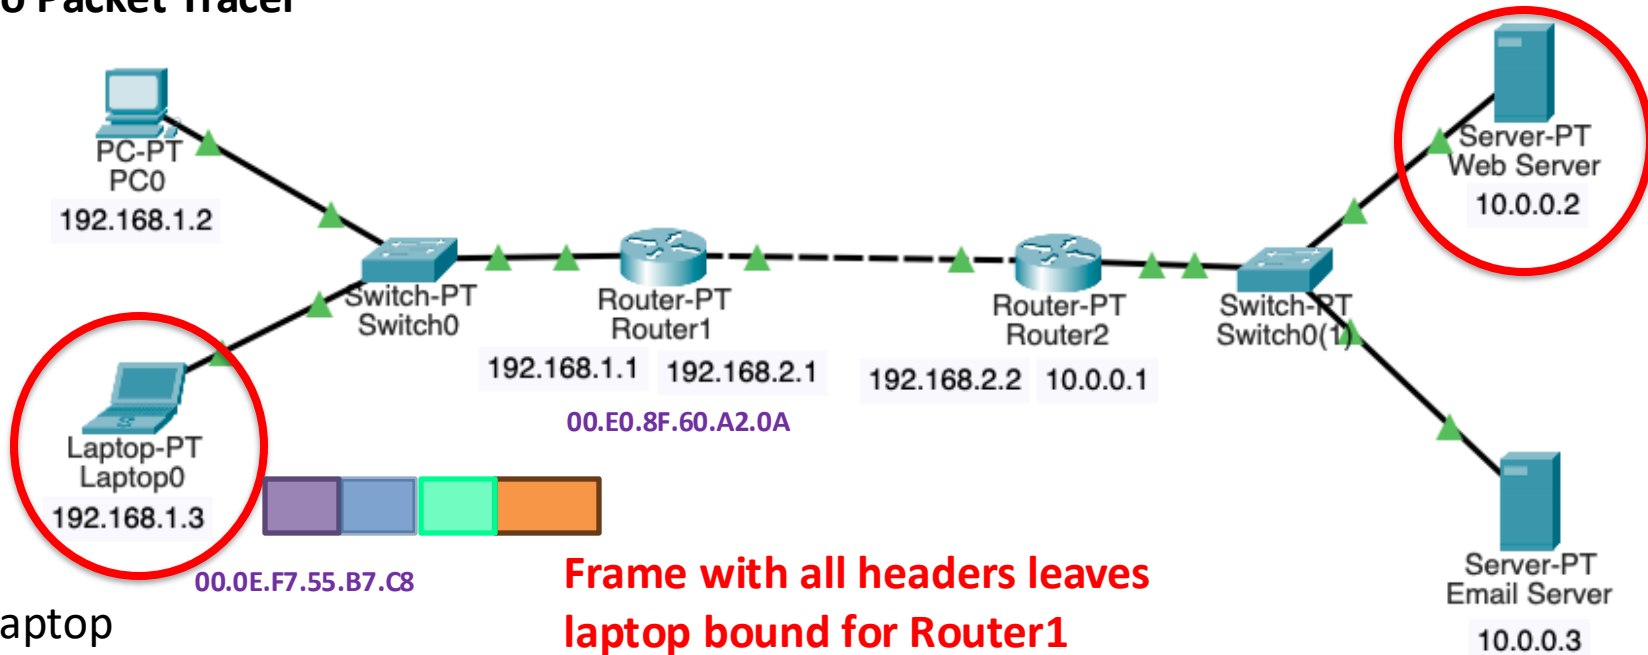
Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Internet largely uses a client/server model where clients make requests to servers

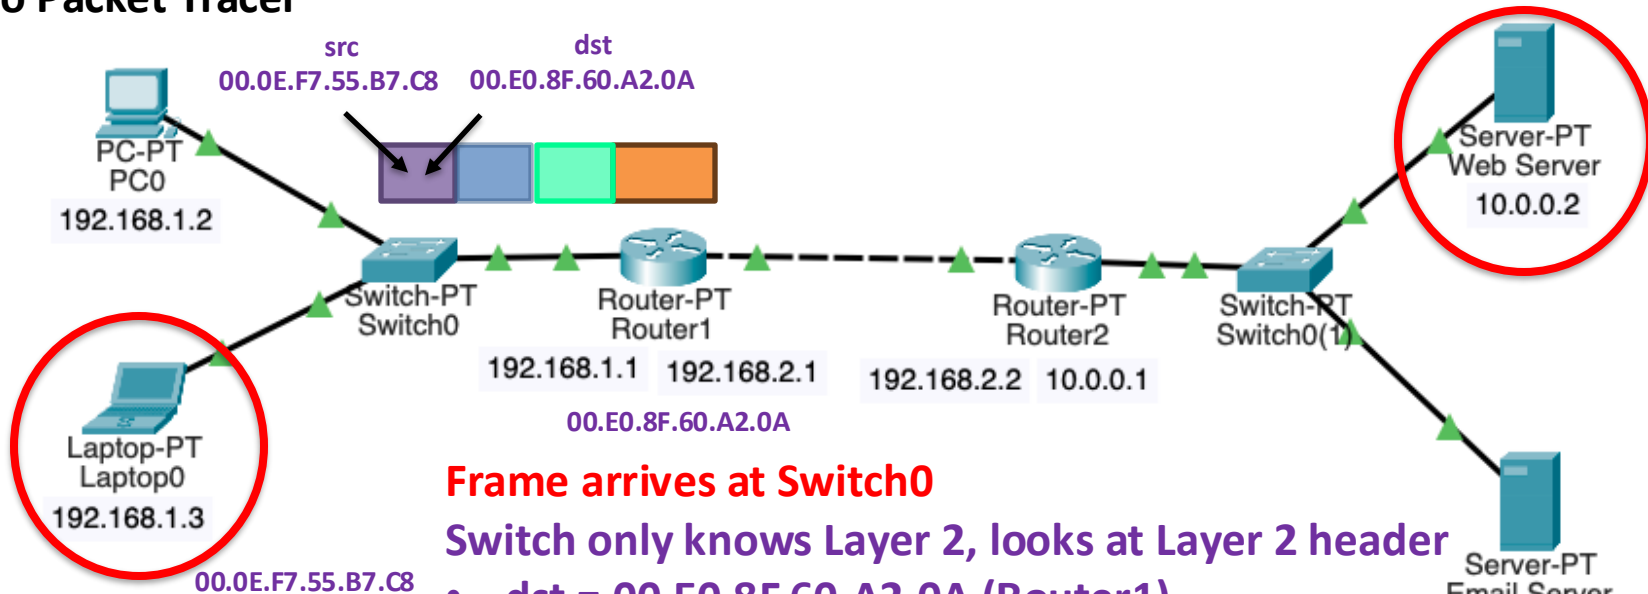
Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Frame arrives at Switch0

Switch only knows Layer 2, looks at Layer 2 header

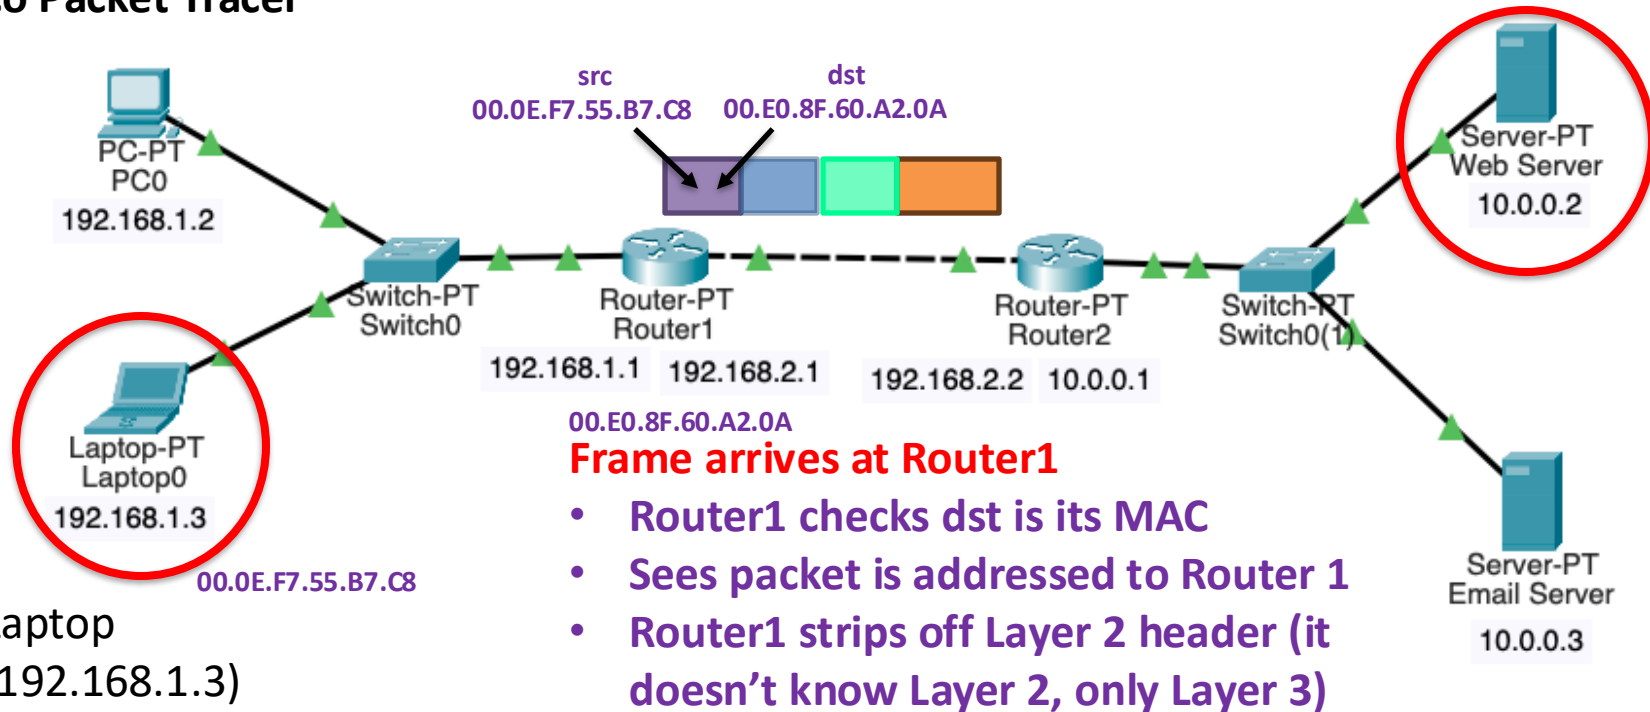
- dst = 00.E0.8F.60.A2.0A (Router1)
- src = 00.0E.F7.55.B7.C8 (Laptop)

Switch knows which physical port dst is plugged into, forwards packet to Router1

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

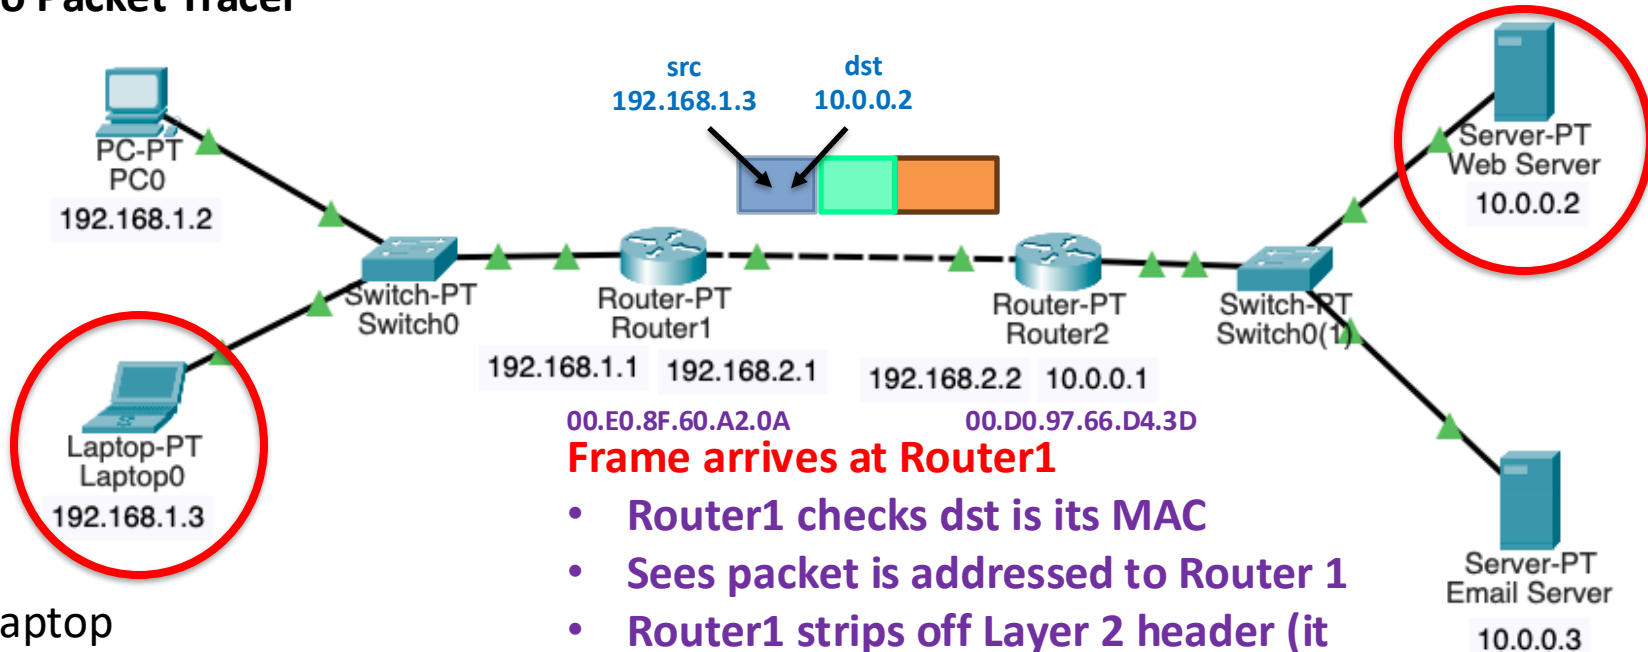
Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



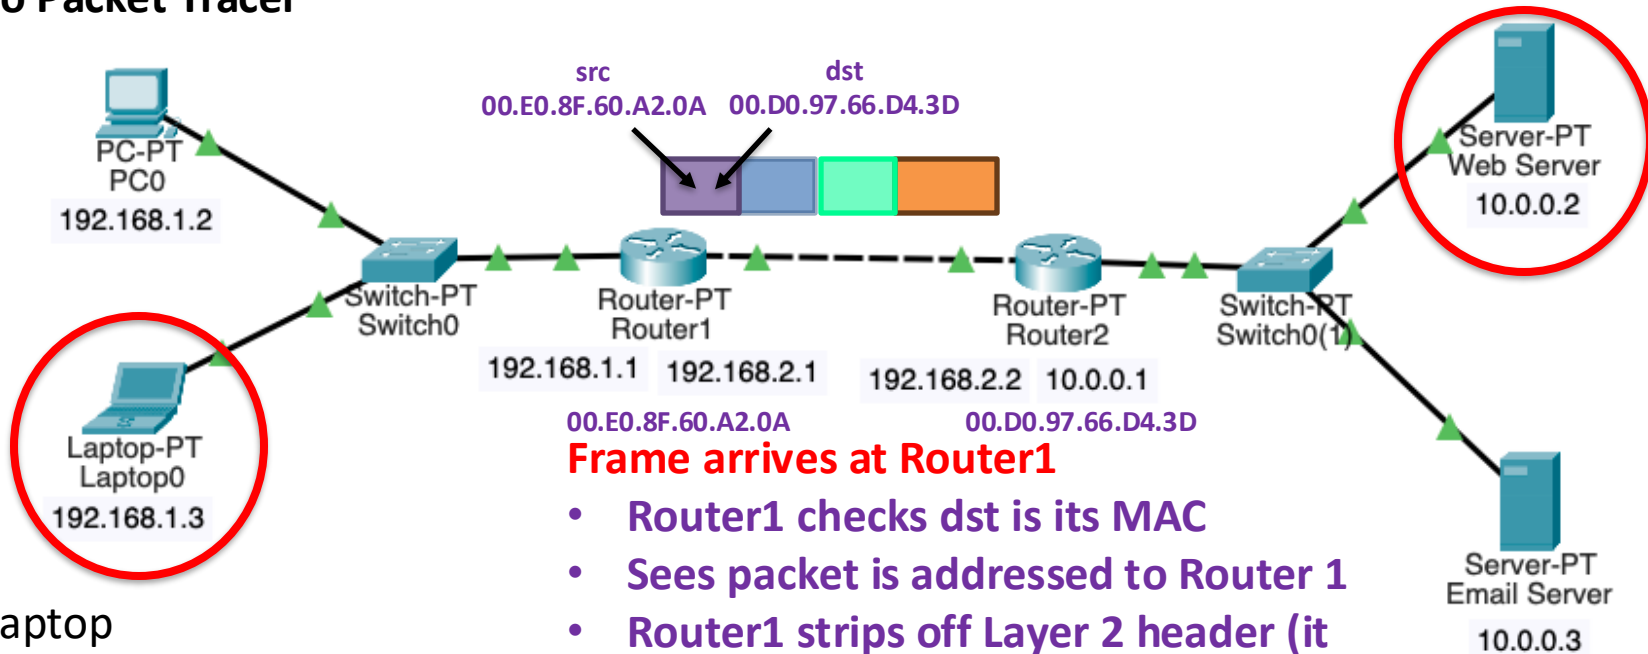
Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Router1

- Router1 checks dst is its MAC
 - Sees packet is addressed to Router 1
 - Router1 strips off Layer 2 header (it doesn't know Layer 2, only Layer 3)
- Router 1 looks at Layer 3 header
- Realizes next stop is Router2 if it wants to send the packet to 10.0.0.2 (we will see how soon)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

**How did Router1 get
the MAC of Router2?
ARP!**

Frame arrives at Router1

- Router1 checks dst is its MAC
- Sees packet is addressed to Router 1
- Router1 strips off Layer 2 header (it doesn't know Layer 2, only Layer 3)

Router 1 looks at Layer 3 header

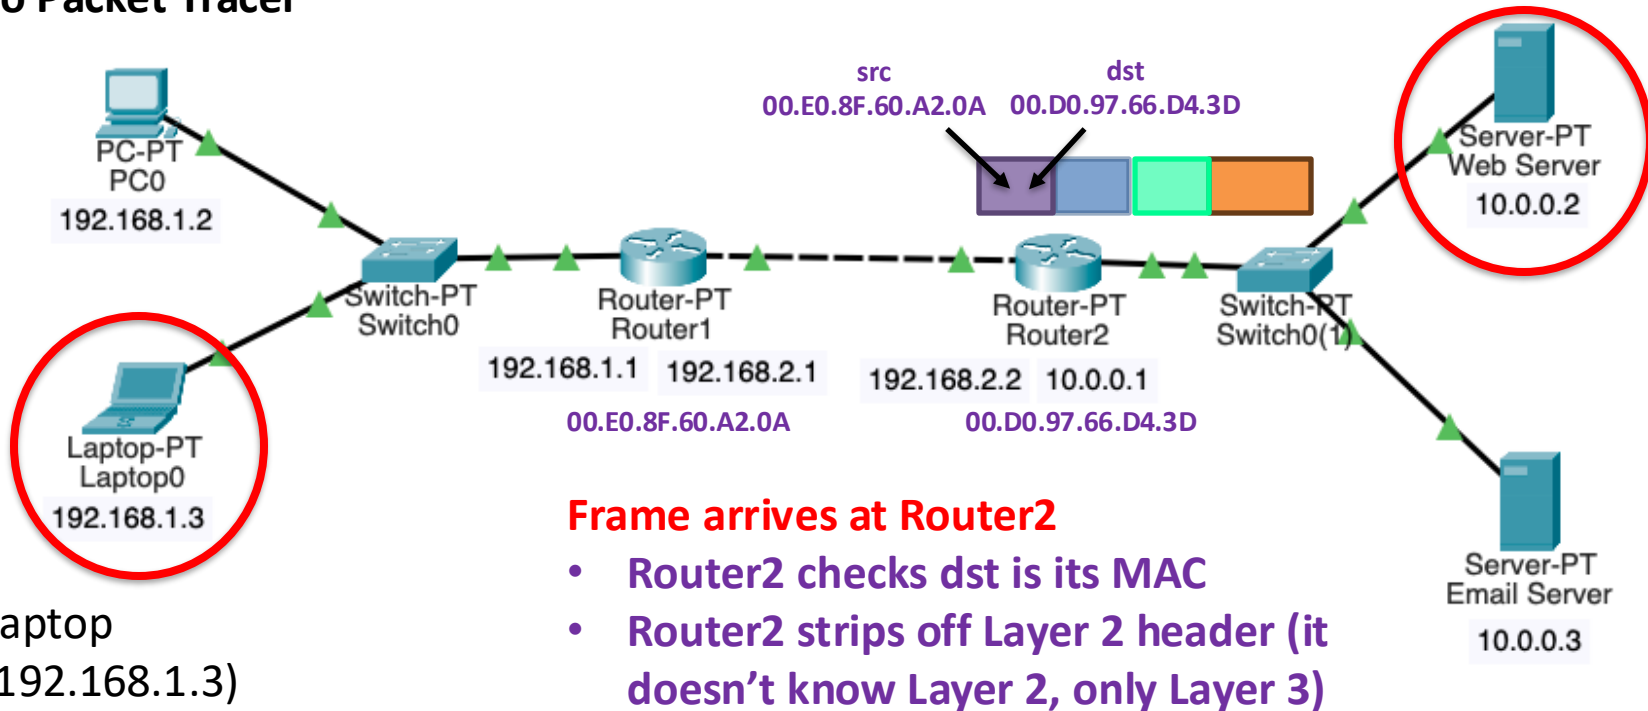
- Realizes next stop is Router2 if it wants to send the packet to 10.0.0.2 (we will see how soon)

Router1 "re-writes" Layer 2 header for Router2

- dst = 00.D0.97.66.D4.3D (Router2)
- src = 00.E0.8F.60.A2.0A (Router1)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



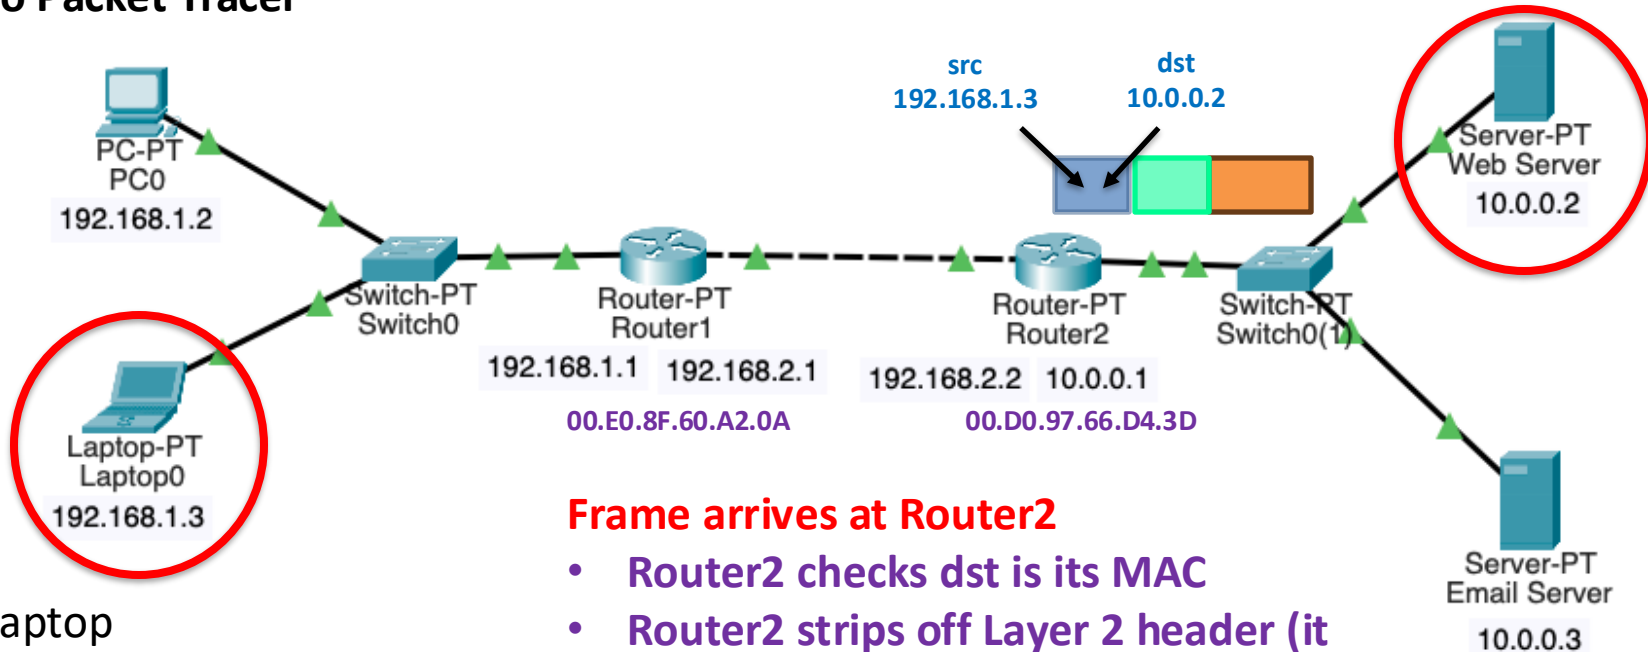
Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Router2

- Router2 checks dst is its MAC
- Router2 strips off Layer 2 header (it doesn't know Layer 2, only Layer 3)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Router2

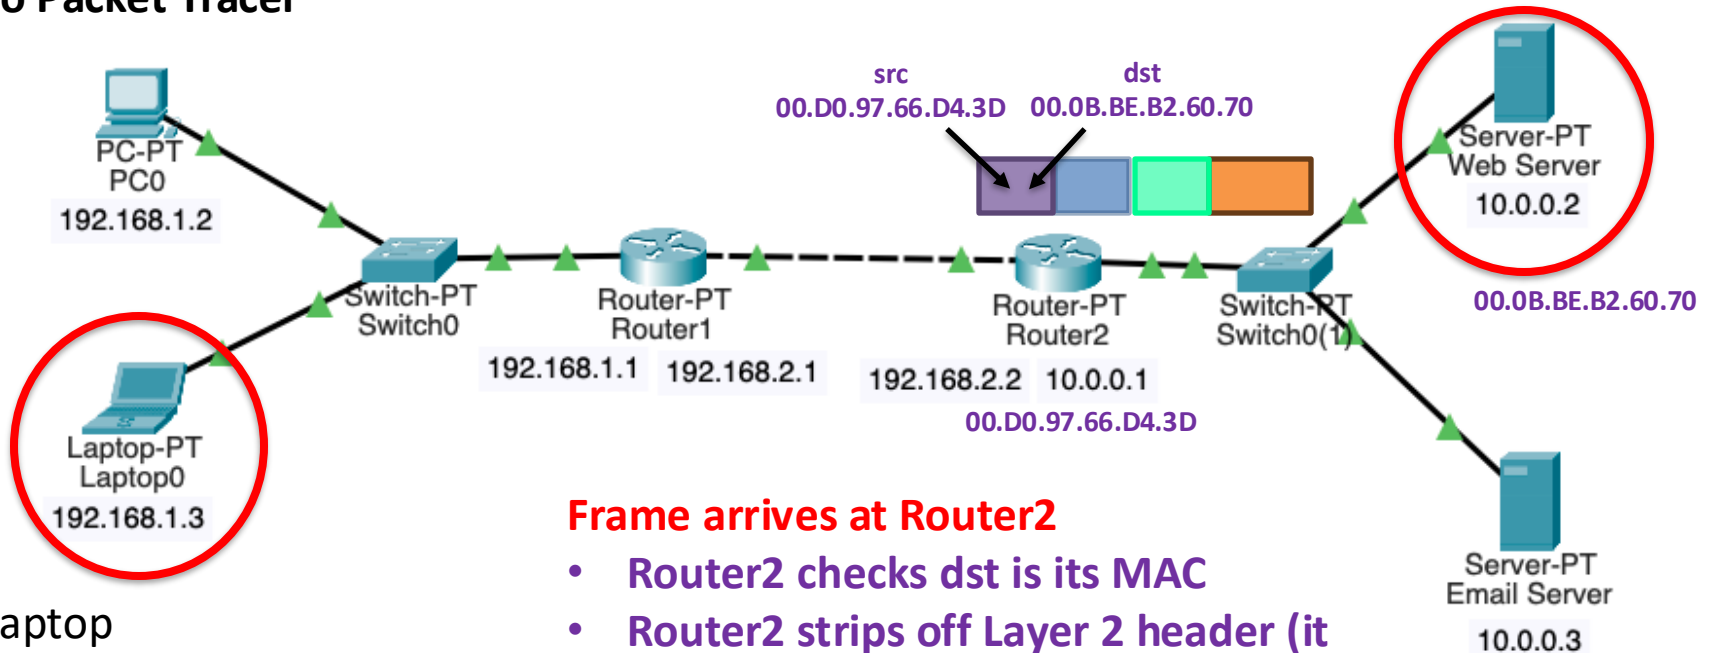
- Router2 checks dst is its MAC
- Router2 strips off Layer 2 header (it doesn't know Layer 2, only Layer 3)

Router 2 looks at Layer 3 header

- dst = 10.0.0.2 (Web Server)
- src = 192.168.1.3 (Laptop)
- Realizes next stop is Web Server (we will see how soon)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Router2

- Router2 checks dst is its MAC
- Router2 strips off Layer 2 header (it doesn't know Layer 2, only Layer 3)

Router 2 looks at Layer 3 header

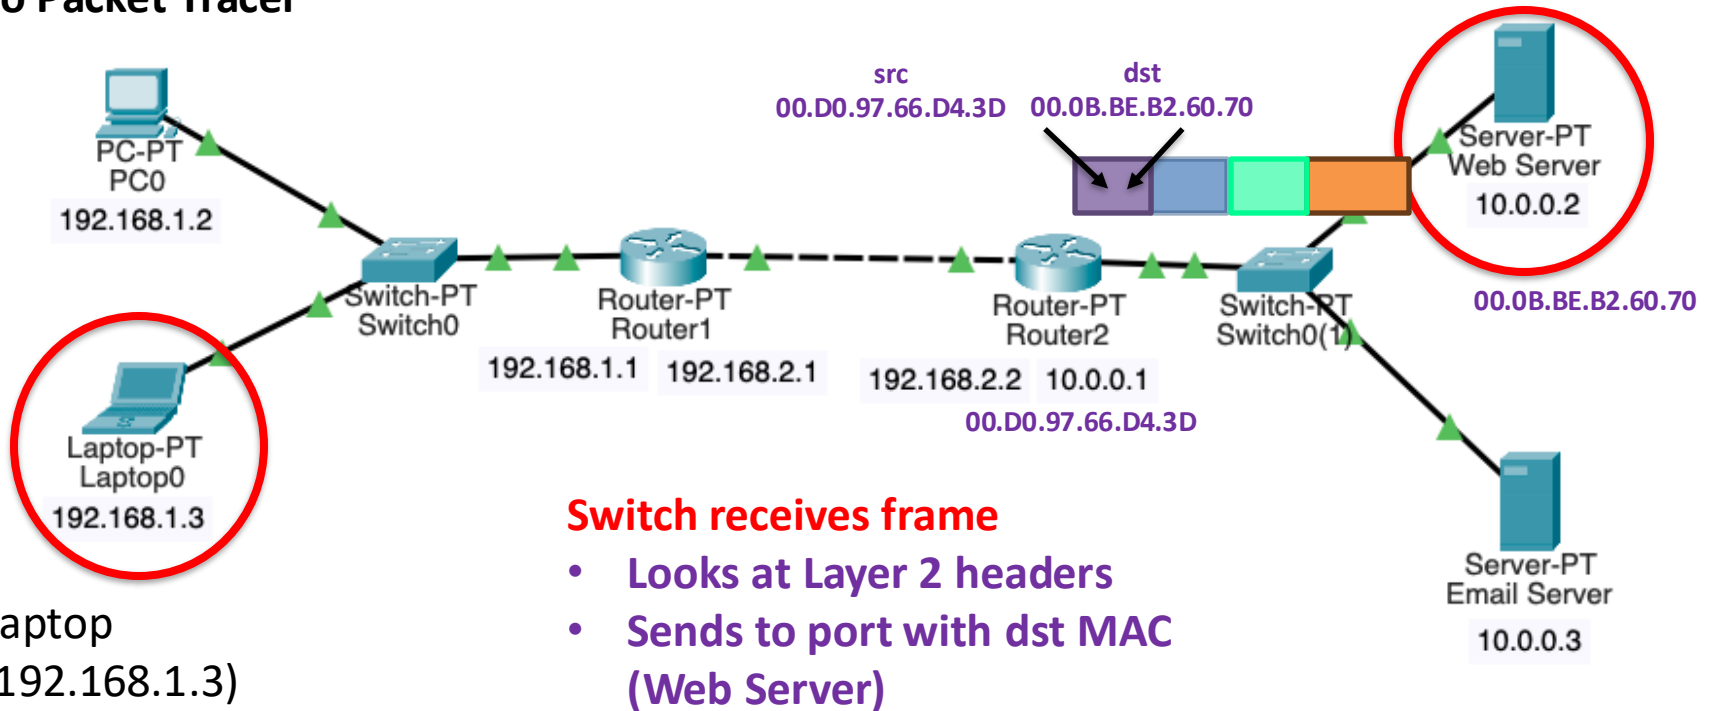
- dst = 10.0.0.2 (Web Server)
- src = 192.168.1.3 (Laptop)
- Realizes next stop is Router2 (we will see how soon)

Router2 "re-writes" Layer 2 header for Web Server

- dst = 00.0B.BE.B2.60.70 (Web server)
- src = 00.D0.97.66.D4.3D (Router 2)

Internet largely uses a client/server model where clients make requests to servers


Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer

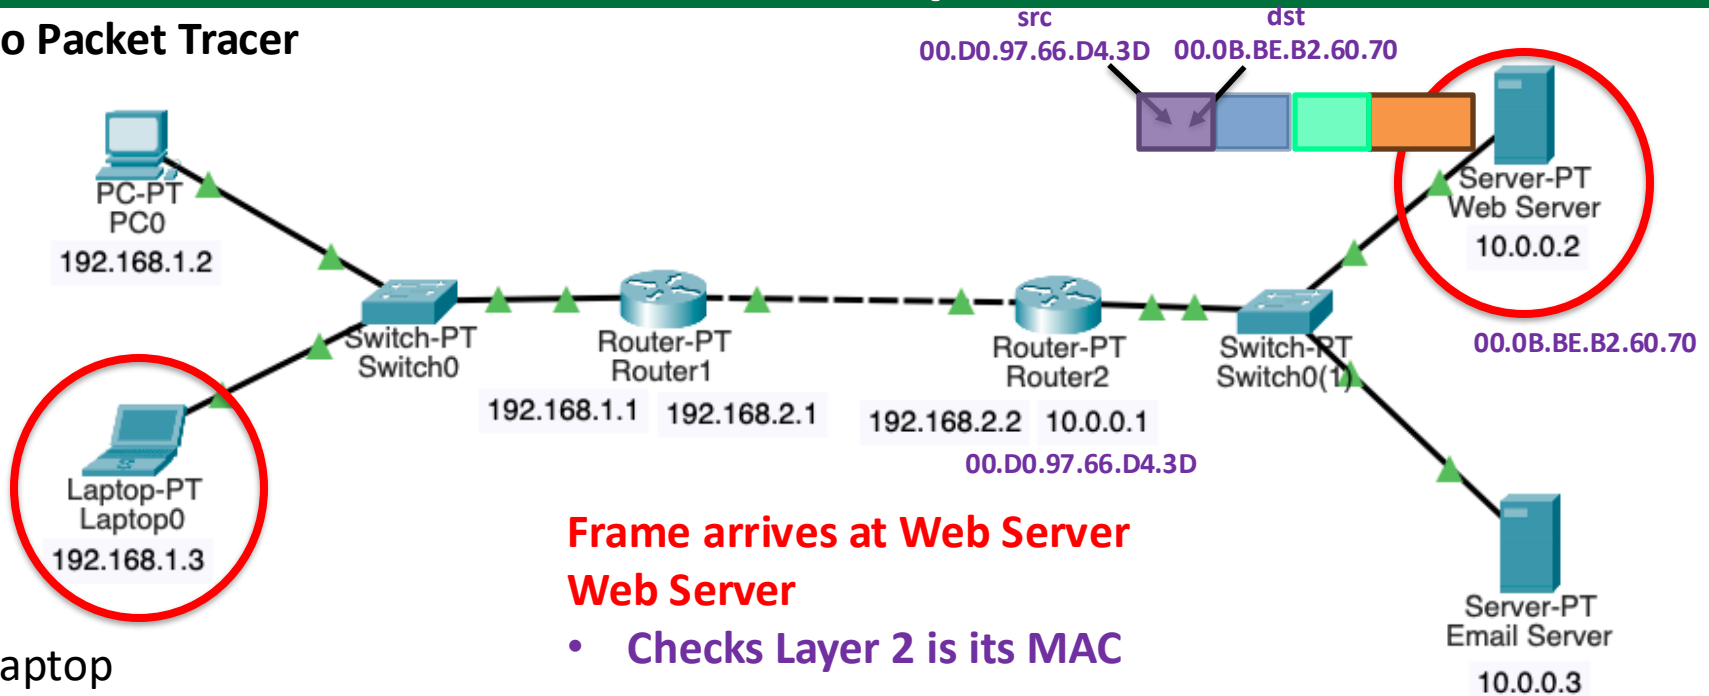


Frame arrives at Web Server

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

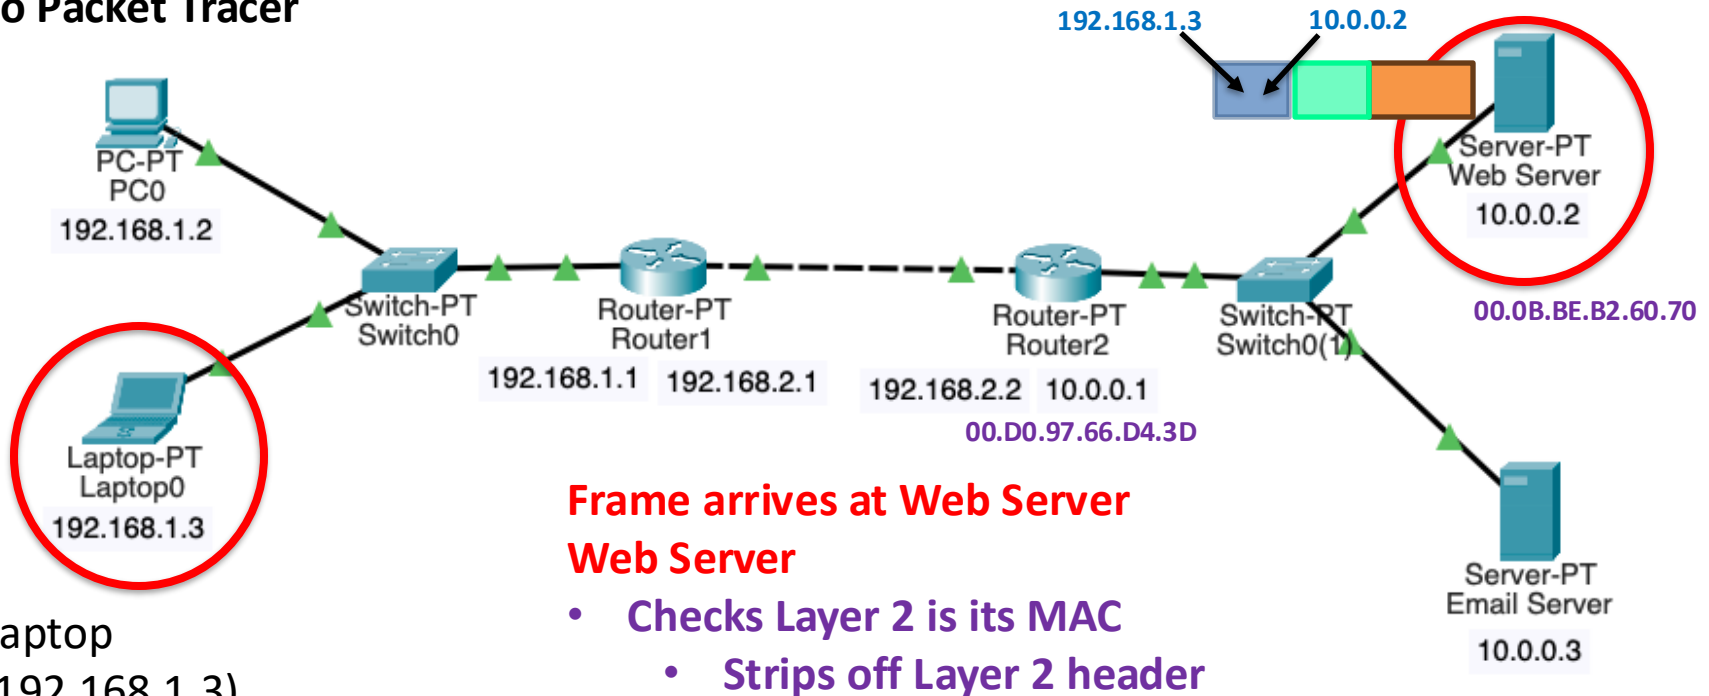
Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers


Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



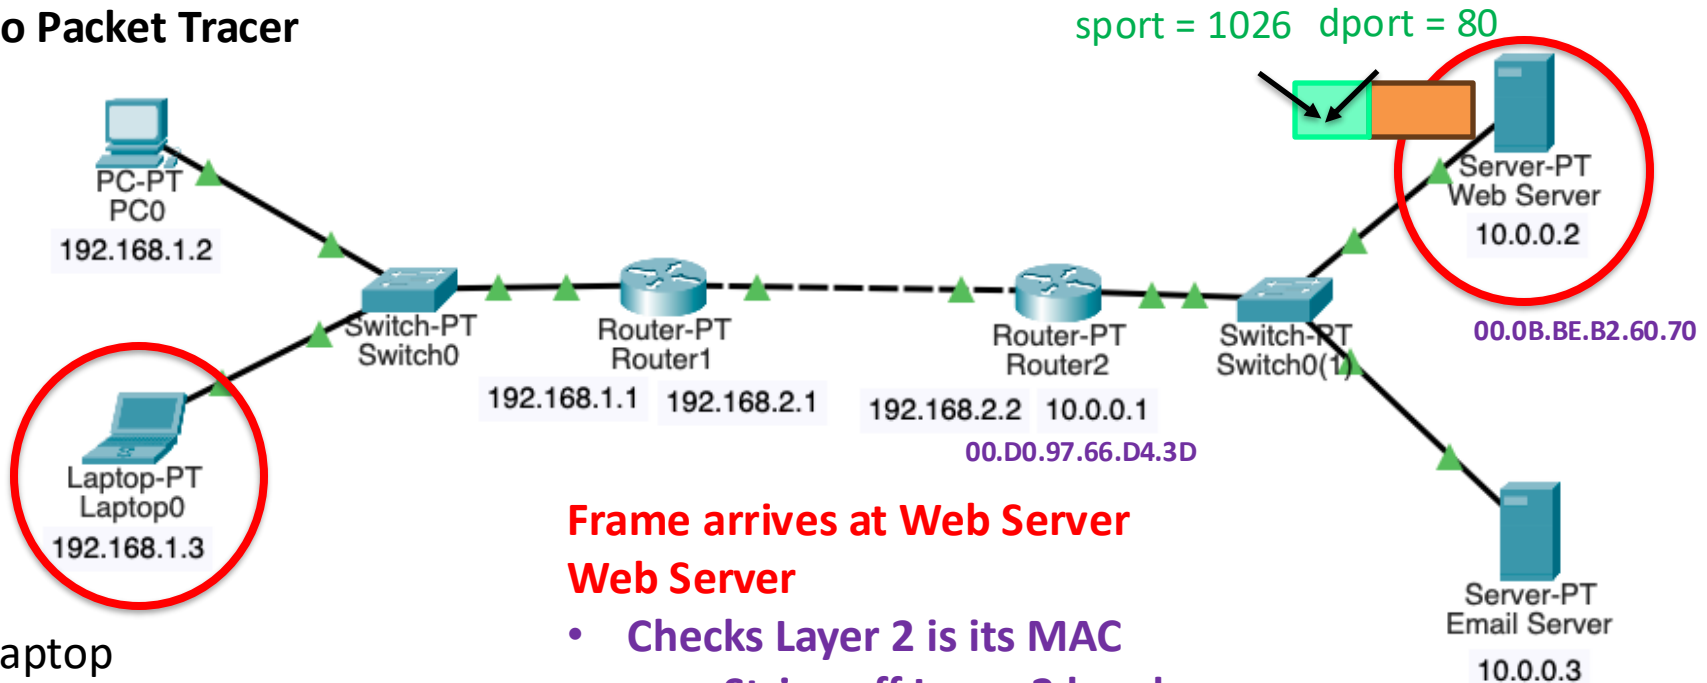
Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Web Server Web Server

- Checks Layer 2 is its MAC
 - Strips off Layer 2 header
- Checks Layer 3 IP address
 - Stores src IP (Laptop at 192.168.1.3)
 - Strips off Layer 3 header

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer




Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Web Server Web Server

- Checks Layer 2 is its MAC
 - Strips off Layer 2 header
- Checks Layer 3 IP address
 - Stores src IP (Laptop at 192.168.1.3)
 - Strips off Layer 3 header
- Checks Layer 4
 - Notes dst port (80=web)
 - Stores sport (1026=Laptop) for return
 - Strips off Layer 4 header

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer




Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Frame arrives at Web Server Web Server

- Checks Layer 2 is its MAC
 - Strips off Layer 2 header
- Checks Layer 3 IP address
 - Stores src IP (Laptop at 192.168.1.3)
 - Strips off Layer 3 header
- Checks Layer 4
 - Notes dst port (80=web)
 - Stores sport (1026=Laptop) for return
 - Strips off Layer 4 header
- Sends data to web app

Internet largely uses a client/server model where clients make requests to servers

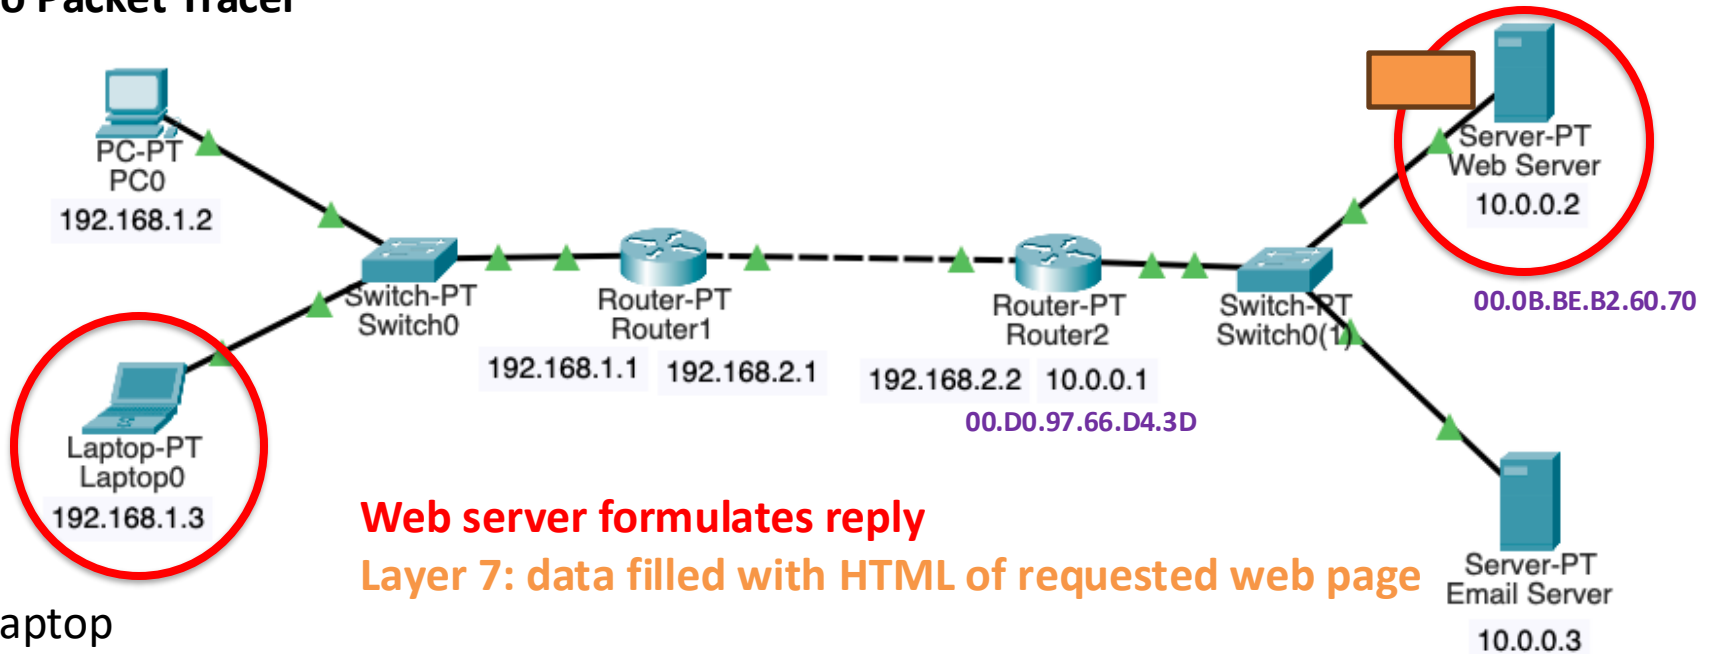
Cisco Packet Tracer



Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



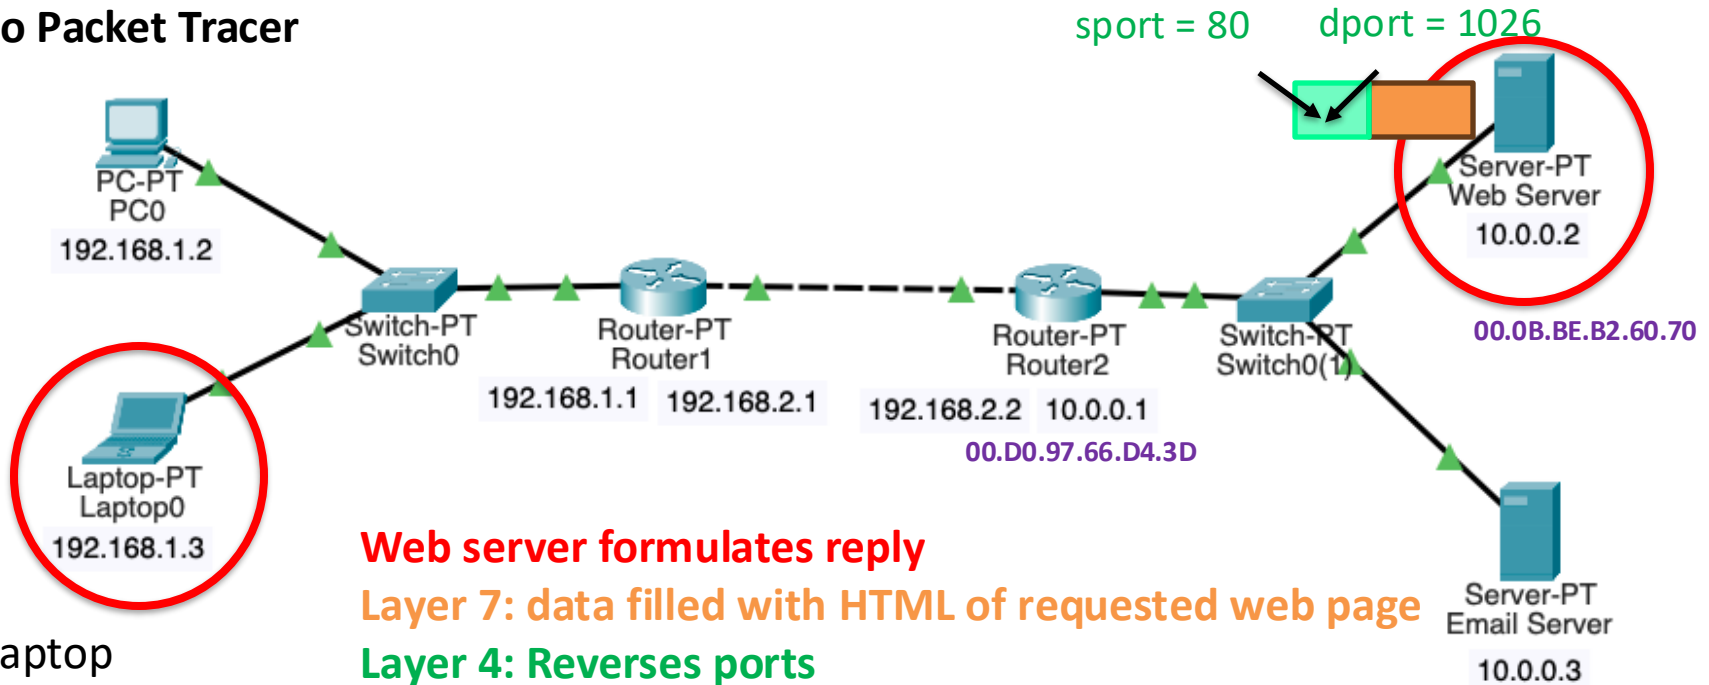
Web server formulates reply

Layer 7: data filled with HTML of requested web page

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model
where clients make requests to servers

Cisco Packet Tracer



Web server formulates reply

Layer 7: data filled with HTML of requested web page

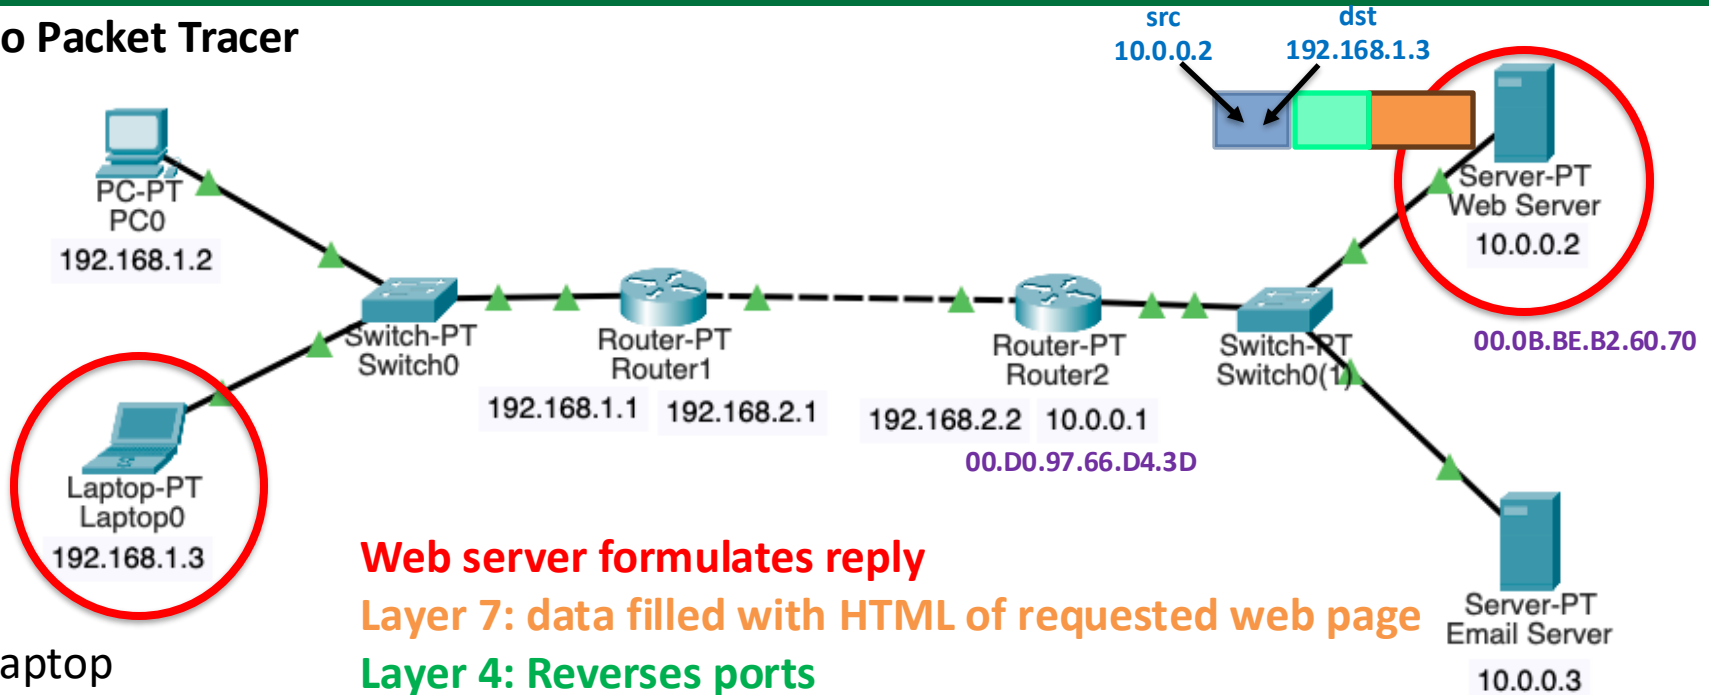
Layer 4: Reverses ports

- sport = 80 (original dport)
- dport = 1026 (original sport)

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Web server formulates reply

Layer 7: data filled with HTML of requested web page

Layer 4: Reverses ports

- **sport = 80 (original dport)**
- **dport = 1026 (original sport)**

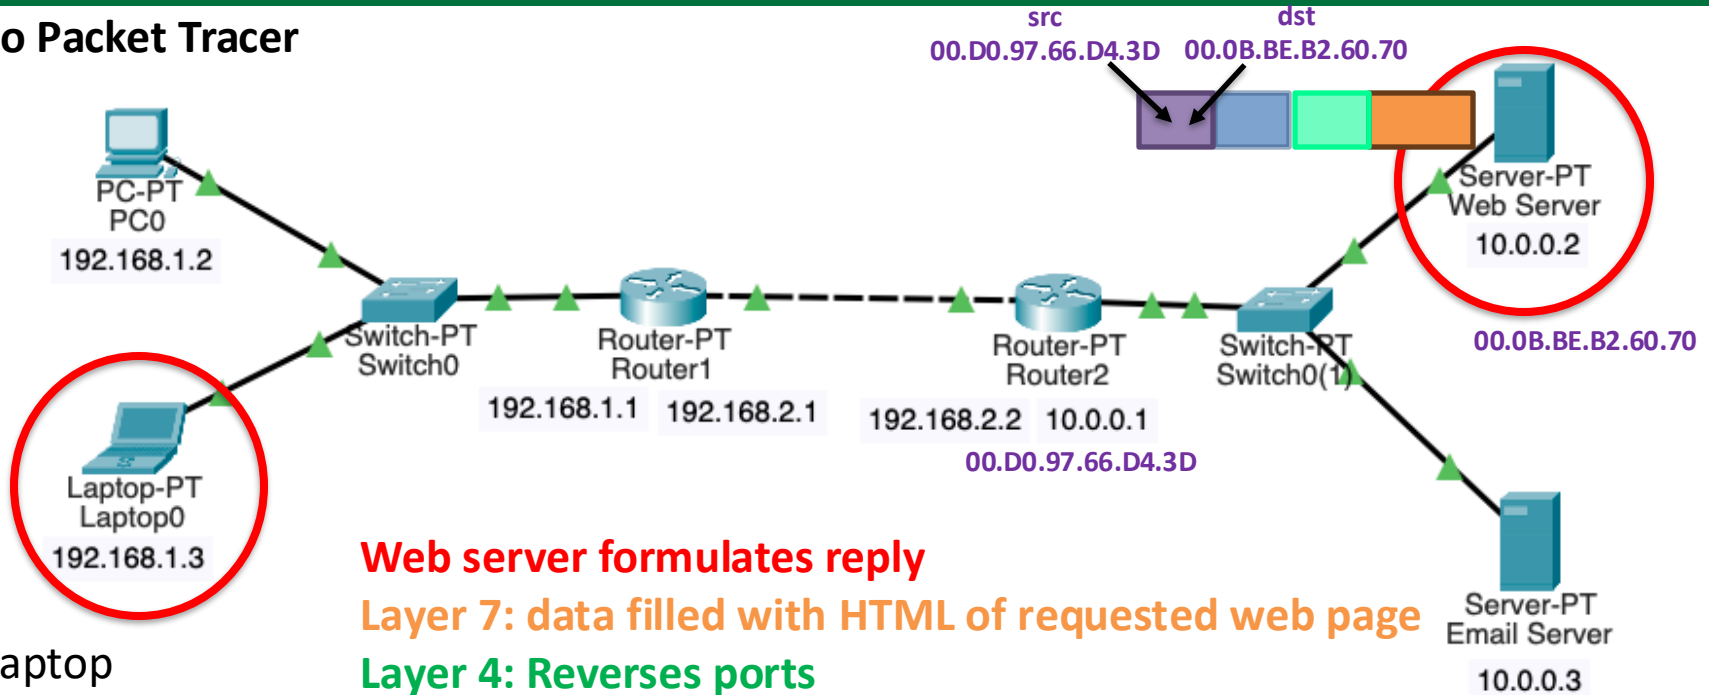
Layer 3: Reverses IP addresses

- **src = 10.0.0.2 (Web Server)**
- **dst = 192.168.1.3 (Laptop)**

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Web server formulates reply

Layer 7: data filled with HTML of requested web page

Layer 4: Reverses ports

- **sport = 80 (original dport)**
- **dport = 1026 (original sport)**

Layer 3: Reverses IP addresses

- **src = 10.0.0.2 (Web Server)**
- **dst = 192.168.1.3 (Laptop)**

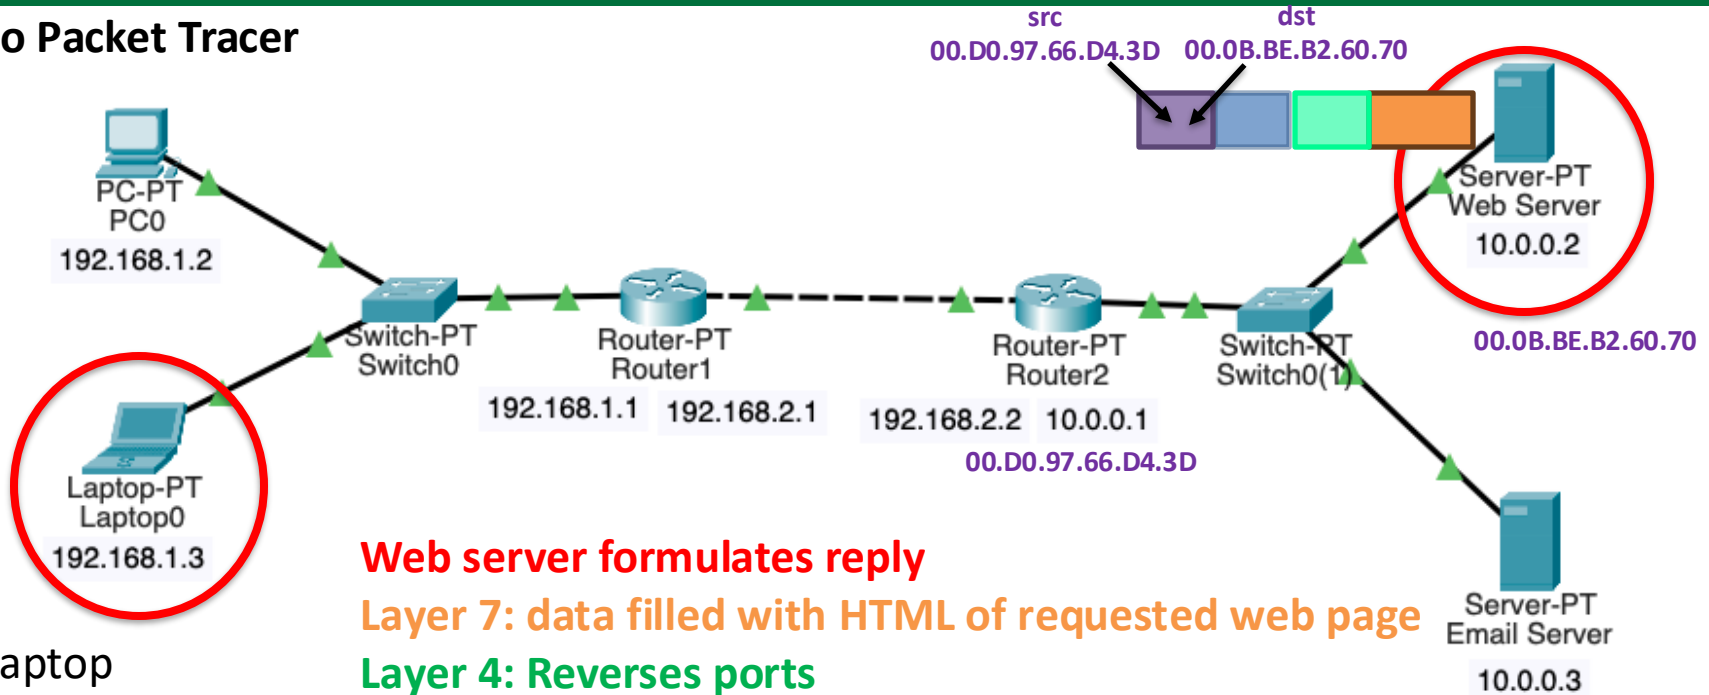
Layer 2: Reverses MAC addresses

- **src = 00.0B.BE.B2.60.70 (Web Server)**
- **dst = 00.D0.97.66.D4.3D (Router2)**

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)

Internet largely uses a client/server model where clients make requests to servers

Cisco Packet Tracer



Web server formulates reply

Layer 7: data filled with HTML of requested web page

Layer 4: Reverses ports

- **sport = 80 (original dport)**
- **dport = 1026 (original sport)**

Layer 3: Reverses IP addresses

- **src = 10.0.0.2 (Web Server)**
- **dst = 192.168.1.3 (Laptop)**

Layer 2: Reverses MAC addresses

- **src = 00.0B.BE.B2.60.70 (Web Server)**
- **dst = 00.D0.97.66.D4.3D (Router2)**

Reply routed back to Laptop

Can the Server get the Laptop's MAC?

No! Router1 stripped it out and overwrote it

Laptop
(192.168.1.3)
makes a request
for a web page
from the Web
Server
(10.0.0.2)


I made one simplification and didn't address NAT

The local network probably uses Network Address Translation (NAT)

- The client's 192.168.1.0/24 address wasn't what was used for routing on the Internet (it is non-routable)
- We will cover NAT soon and why we use it (short answer: we ran out of IP addresses)



AA:BB:CC:DD:EE:FF
192.168.60.5



Agenda

1. Review: network layers




2. Sockets

3. Scapy

4. Exercises

Socket: an endpoint for two-way comms between programs running on a network




Sockets are defined by:

- IP address
- Port
- Protocol (TCP or UDP)

Sockets are an interface for an application to send or receive data from the network

Typically, the OS handles packet construction



**Each NIC has unique MAC address
Drops packets not addresses to itself
(unless promiscuous or monitor mode)**

Application has data that it wants to send

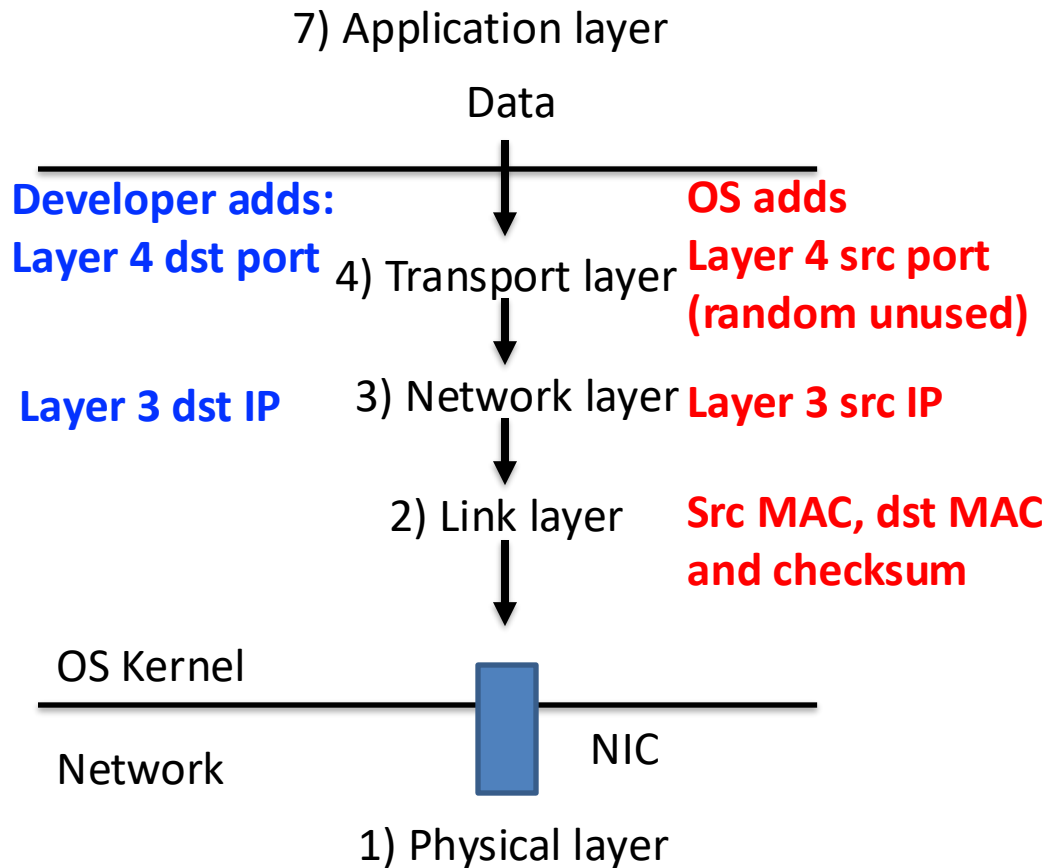
Application passes data to OS via a socket

OS generates a packet with Transport, Network, and Link Layer headers

Packet sent out over network via NIC

**How does the application hand data to OS?
OS provides socket API system calls
Recall sockets from CS10 and CS50**

Applications must provide protocol and destination information



Application developer must decide which transport protocol to use:

- UDP – faster, but delivery not guaranteed
- TCP – slower but delivery guaranteed

Developer must also provide destination information

- IP address (Layer 3)
- Port (Layer 4)

OS adds source port and IP, MAC and checksum

Sending a packet using a socket is not difficult with Python

send_udp.py

We will use a socket (as with CS10 and CS50)

```
import socket
```

```
dest_addr = "127.0.0.1"
```

```
port = 9090
```

```
msg = b'Hello world!'
```

Application's data (in bytes) to send

```
if __name__ == '__main__':
```

```
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
    sock.sendto(msg,(dest_addr,port))
```

```
    sock.close()
```

Socket will use IPv4

This Socket will use UDP protocol
Set to socket.SOCK_STREAM for TCP

Send 'Hello world'

To destination IP address
(localhost = 127.0.0.1)

Using port 9090

OS handles adding Transport, Network, and Link Layer headers to applications data
OS sends headers and data over Physical Layer 1 via NIC

Sending a packet using a socket is not difficult with Python

send_udp.py

Python hides a lot of complexity!

```
import socket
```

```
dest_addr = "127.0.0.1"
```

```
port = 9090
```

```
msg = b'Hello world!'
```

```
if __name__ == '__main__':
```

```
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
    sock.sendto(msg,(dest_addr,port))
```

Note: we do not set a source port on the client side

OS picks a random unused port for us

We do set a port on the server side

Client needs to know what port the server is listening on

Run program

Terminal 1

```
$ python3 send_udp.py
```

Listen for UDP in a terminal on port 9090 using netcat (nc)

Terminal 2

```
$ nc -l -u 9090  
Hello world!
```

netcat receives UDP message when sent

Sending a packet via a Socket in C is a little more involved, but gives you more control

send_udp.c

```
int main() {  
    struct sockaddr_in dest_info;  
    char *dest_addr = "127.0.0.1";  
    int port = 9090;  
    char *data = "Hello World (in C!)\n";  
  
    //Create network socket  
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);  
  
    //Provide needed data  
    memset((char *) &dest_info, 0, sizeof(dest_info));  
    dest_info.sin_family = AF_INET;  
    dest_info.sin_addr.s_addr = inet_addr(dest_addr);  
    dest_info.sin_port = htons(port);  
  
    //send packet  
    sendto(sock, data, strlen(data), 0, (struct sockaddr *) &dest_info,  
          sizeof(dest_info));  
  
    close(sock);  
}
```

Python built on
top of C libraries

Create socket for
UDP



Use IPv4



Set destination info



Send



sizeof(dest_info));

Sending a packet via a Socket in C is a little more involved, but gives you more control

send_udp.c

```
int main() {  
    struct sockaddr_in dest_info;  
    char *dest_addr = "127.0.0.1";  
    int port = 9090;  
    char *data = "Hello World (in C!)\n";  
  
    //Create network socket  
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);  
  
    //Provide needed data  
    memset((char *) &dest_info, 0, sizeof(dest_info));  
    dest_info.sin_family = AF_INET;  
    dest_info.sin_addr.s_addr = inet_addr(dest_addr);  
    dest_info.sin_port = htons(port);  
  
    //send packet  
    sendto(sock, data, strlen(data), 0, (struct sockaddr *) &dest_info,  
    clo
```

Python built on
top of C libraries

Create socket for
UDP

Use IPv4

Set destination info

Send

Listen for UDP
in a terminal

netcat receives
UDP message
when sent

Terminal 1

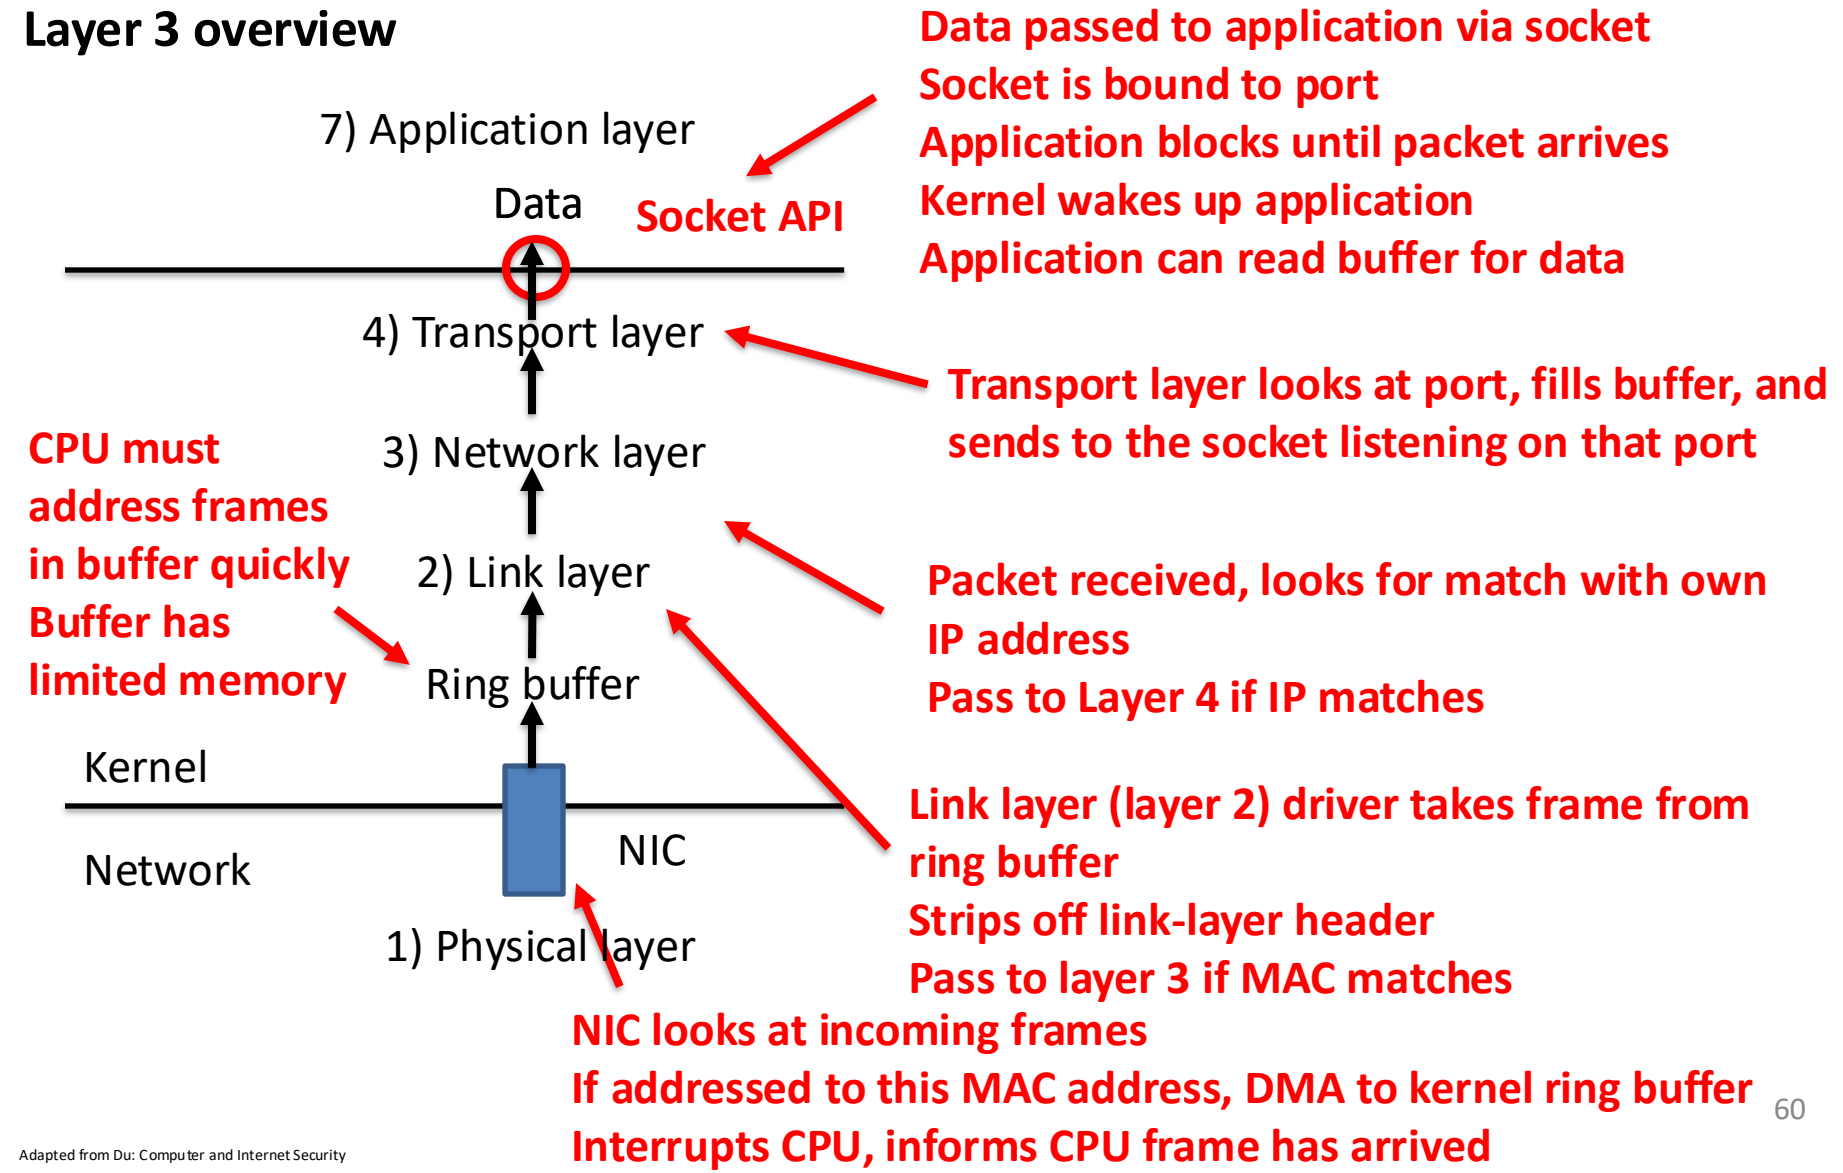
```
$ gcc gcc send_udp.c -o send  
$ ./send
```

Terminal 2

```
$ nc -luv 9090  
Hello World (in C)!
```

Receiving a packet via a Socket is the reverse of sending a packet

Layer 3 overview



UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes
```

```
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces
```

```
port = 9090 #listen on this port
```

```
if __name__ == '__main__':
```

```
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
    sock.bind((ip_addr,port))
```

Set up socket for UDP as before



```
while (True):
```

```
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)
```

```
    print(f"{ip}:{port} {msg.decode('utf-8')}")
```

UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

Computer might have multiple NICs

ip_addr of zeros means accept packets from any of them

```
MAX_SIZE = 1500 #max message size in bytes
```

```
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces
```

```
port = 9090 #listen on this port
```

Set up socket for UDP as before

```
if __name__ == '__main__':
```

```
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
    sock.bind((ip_addr, port))
```

Bind socket to listen on ip_addr and port
Note: we pick a port on the server (receiver) side

```
while (True):
```

```
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)
```

```
    print(f"{ip}:{port} {msg.decode('utf-8')}")
```

UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes  
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces  
port = 9090 #listen on this port
```

```
if __name__ == '__main__':  
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)  
    sock.bind((ip_addr, port))
```

Loop forever

```
while (True):  
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)  
    print(f'{ip}:{port} {msg.decode('utf-8')}')
```

Block until packet arrives
recvfrom gets data
here print client ip, port,
and msg

Use recvfrom for UDP
Use recv for TCP

UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes  
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces  
port = 9090 #listen on this port
```

```
if __name__ == '__main__':  
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)  
    sock.bind((ip_addr,port))
```

```
while (True):  
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)  
    print(f'{ip}:{port} {msg.decode('utf-8')}')
```

Terminal 1 (receiver)

```
$ python3 receive_udp.py
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090
```


UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes  
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces  
port = 9090 #listen on this port
```

```
if __name__ == '__main__':  
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)  
    sock.bind((ip_addr,port))
```

```
while (True):  
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)  
    print(f'{ip}:{port} {msg.decode('utf-8')}')
```

Sender types msg

Terminal 1 (receiver)

```
$ python3 receive_udp.py
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090  
hello
```

UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes  
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces  
port = 9090 #listen on this port
```

```
if __name__ == '__main__':  
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)  
    sock.bind((ip_addr, port))
```

```
    while (True):  
        msg, (ip, port) = sock.recvfrom(MAX_SIZE)  
        print(f'{ip}:{port} {msg.decode('utf-8')}')  
        # Receiver gets msg,  
        # ip addr, and port
```

Sender types msg

Terminal 1 (receiver)

```
$ python3 receive_udp.py  
127.0.0.1:64369 hello
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090  
hello
```

UDP receiver in Python using a Socket

receive_udp.py

We did not pick a source port on the sender (Terminal 2)

OS chooses source port randomly

When this is run, see source port other than 9090 (it was 64369 here)

On receiver, OS can sort out replies if multiple instances of application are running at the same time based on port

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes
```

```
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces
```

```
port = 9090 #listen on this port
```

```
if __name__ == '__main__':
```

```
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
    sock.bind((ip_addr,port))
```

```
    while (True):
```

```
        msg, (ip, port) = sock.recvfrom(MAX_SIZE)
```

```
        print(f'{ip}:{port} {msg.decode("utf-8")}')
```

Receiver gets msg,
ip addr, and port

Sender types msg

Terminal 1 (receiver)

```
$ python3 receive_udp.py
```

```
127.0.0.1:64369 hello
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090
```

```
hello
```

UDP receiver in Python using a Socket

receive_udp.py

```
import socket
```

```
MAX_SIZE = 1500 #max message size in bytes  
ip_addr = "0.0.0.0" #0.0.0.0 means bind to all interfaces  
port = 9090 #listen on this port
```

```
if __name__ == '__main__':  
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)  
    sock.bind((ip_addr, port))
```

```
while (True):  
    msg, (ip, port) = sock.recvfrom(MAX_SIZE)  
    print(f'{ip}:{port} {msg.decode('utf-8')}')
```

Run with two senders on the same computer (localhost)
Each sender got different source port
Receiver can tell them apart

Terminal 1 (receiver)

```
$ python3 receive_udp.py  
127.0.0.1:53072 hello  
127.0.0.1:52815 sup?
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090  
hello
```

Terminal 3 (sender)

```
$ nc -u 127.0.0.1 9090  
sup?
```

UDP receiver in C using Sockets is not much different from Python

receive_udp.c

```
const int MAX_SIZE = 1500;
const int port = 9090;
```

```
int main() {
    struct sockaddr_in server;
    struct sockaddr_in client;
    unsigned int clientlen;
    char buf[MAX_SIZE];
    char ip[INET_ADDRSTRLEN];

    // Create the socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    memset((char *) &server, 0, sizeof(server));
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(port);

    if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0) {
        printf("Binding error!");
        return(EXIT_FAILURE);
    }

    // Getting captured packets
    while (1) {
        bzero(buf, MAX_SIZE);
        recvfrom(sock, buf, MAX_SIZE-1, 0, (struct sockaddr *) &client, &clientlen);
        //get ip address as string
        inet_ntop(AF_INET, &(client.sin_addr), ip, INET_ADDRSTRLEN);
        printf("%s:%i %s\n", ip, client.sin_port, buf);
    }
    close(sock);
    return(EXIT_SUCCESS);
}
```

Create socket named sock:

- **AF_INET = IPv4**
- **SOCK_DGRAM = UDP**

Create server and fill memory with zeros

Set server:

- **AF_INET = IPv4**
- **INADDR_ANY = use any interface like Python 0.0.0.0**
- **Port = 9090**

Bind socket to server

**Loop
forever**

**Msg buffer and client info filled
when packets arrive**

UDP receiver in C using Sockets is not much different from Python

receive_udp.c

```
const int MAX_SIZE = 1500;
const int port = 9090;

int main() {
    struct sockaddr_in server;
    struct sockaddr_in client;
    unsigned int clientlen;
    char buf[MAX_SIZE];
    char ip[INET_ADDRSTRLEN];

    // Create the socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    memset((char *) &server, 0, sizeof(server));
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(port);

    if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0) {
        printf("Binding error!");
        return(EXIT_FAILURE);
    }

    // Getting captured packets
    while (1) {
        bzero(buf, MAX_SIZE);
        recvfrom(sock, buf, MAX_SIZE-1, 0, (struct sockaddr *) &client, &clientlen);
        //get ip address
        inet_ntop(AF_INET, &(client.sin_addr), ip, INET_ADDRSTRLEN);
        printf("%s:%i %s\n", ip, client.sin_port, buf);
    }
    close(sock);
    return(EXIT_SUCCESS);
}
```

Terminal 1 (receiver)

```
$ gcc receive_udp.c -o receive
./receive
```

UDP receiver in C using Sockets is not much different from Python

receive_udp.c

```
const int MAX_SIZE = 1500;
const int port = 9090;

int main() {
    struct sockaddr_in server;
    struct sockaddr_in client;
    unsigned int clientlen;
    char buf[MAX_SIZE];
    char ip[INET_ADDRSTRLEN];

    // Create the socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    memset((char *) &server, 0, sizeof(server));
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(port);

    if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0) {
        printf("Binding error!");
        return(EXIT_FAILURE);
    }

    // Getting captured packets
    while (1) {
        bzero(buf, MAX_SIZE);
        recvfrom(sock, buf, MAX_SIZE-1, 0, (struct sockaddr *) &client, &clientlen);
        //get ip address
        inet_ntop(AF_INET, &(client.sin_addr), ip, INET_ADDRSTRLEN);
        printf("%s:%i %s\n", ip, client.sin_port, buf);
    }
    close(sock);
    return(EXIT_SUCCESS);
}
```

Terminal 1 (receiver)

```
$ gcc receive_udp.c -o receive
./receive
```

Terminal 2 (sender)

```
$ nc -u 127.0.0.1 9090
hello
```

UDP receiver in C using Sockets is not much different from Python

receive_udp.c

```
const int MAX_SIZE = 1500;
const int port = 9090;

int main() {
    struct sockaddr_in server;
    struct sockaddr_in client;
    unsigned int clientlen;
    char buf[MAX_SIZE];
    char ip[INET_ADDRSTRLEN];

    // Create the socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);
    memset((char *) &server, 0, sizeof(server));
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(port);

    if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0) {
        printf("Binding error!");
        return(EXIT_FAILURE);
    }

    // Getting captured packets
    while (1) {
        bzero(buf, MAX_SIZE);
        recvfrom(sock, buf, MAX_SIZE-1, 0, (struct sockaddr *) &client, &clientlen);
        //get ip address
        inet_ntop(AF_INET, &(client.sin_addr), ip, INET_ADDRSTRLEN);
        printf("%s:%i %s\n", ip, client.sin_port, buf);
    }
    close(sock);
    return(EXIT_SUCCESS);
}
```

Terminal 1 (receiver)

```
$ gcc receive_udp.c -o receive
./receive
127.0.0.1:14809 hello
```

Terminal 2 (sender)

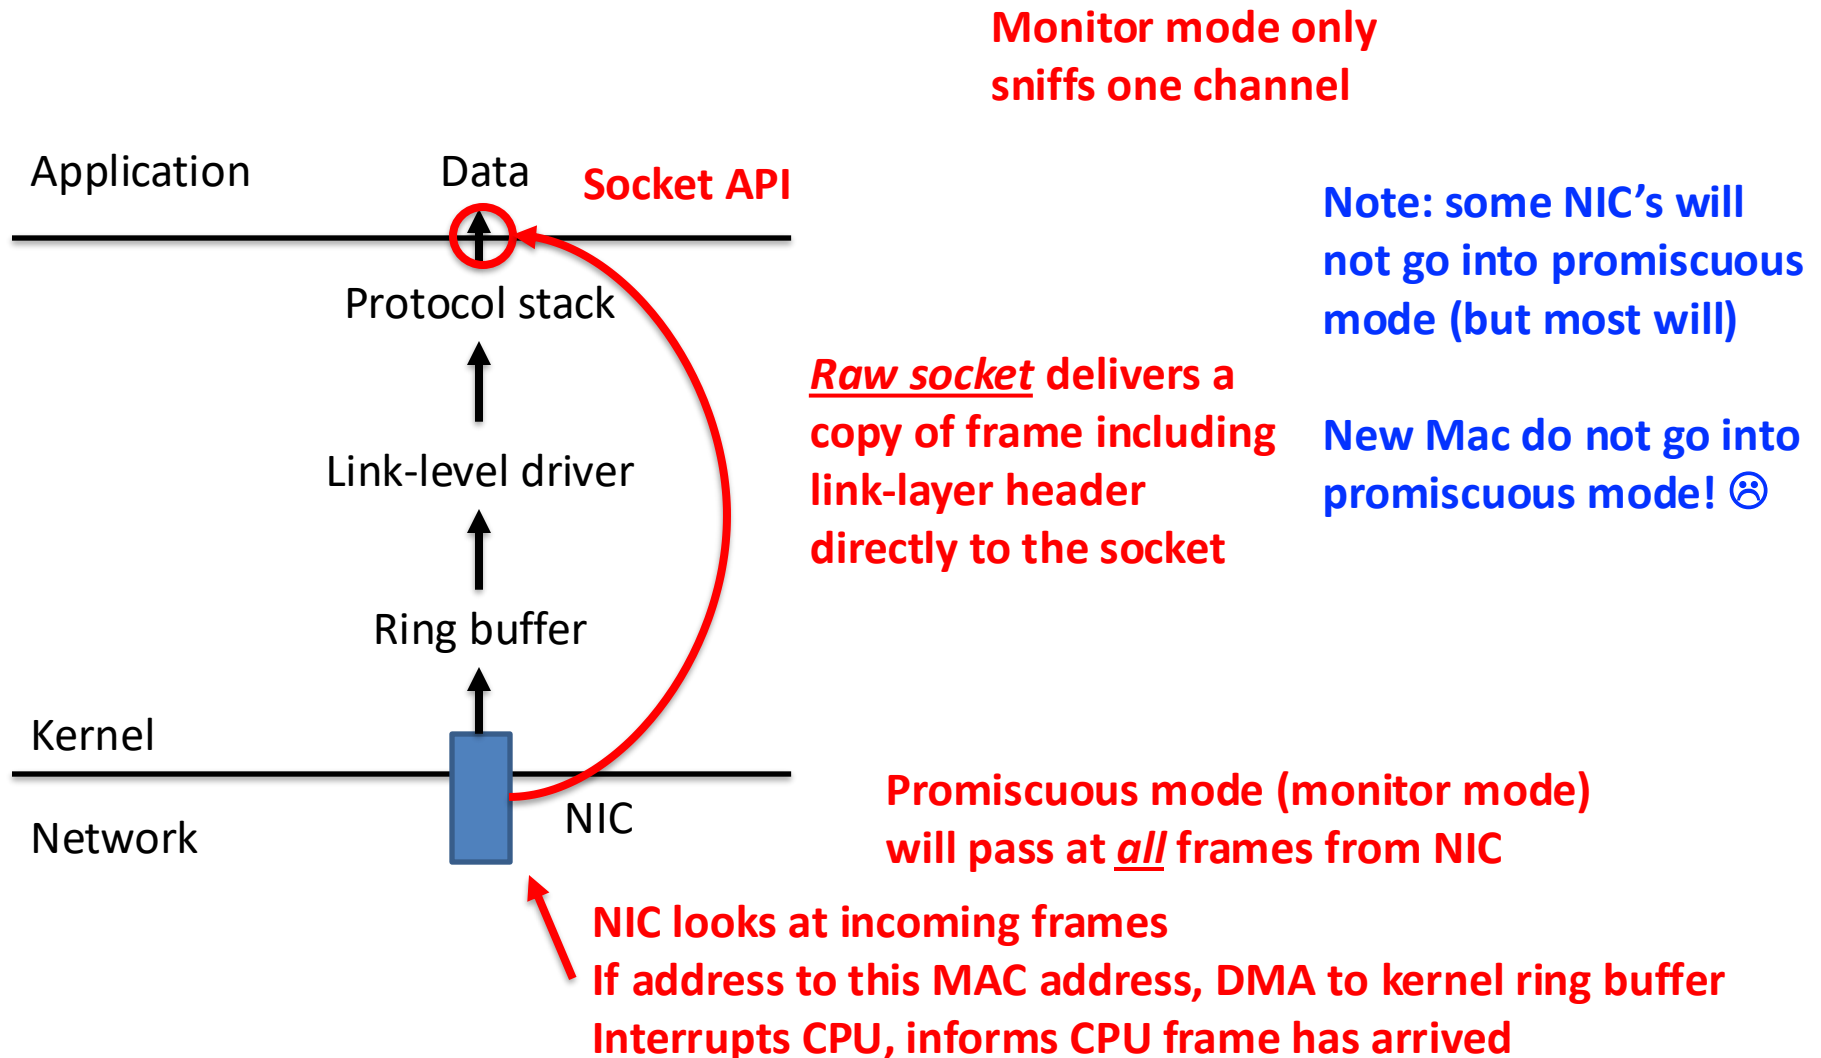
```
$ nc -u 127.0.0.1 9090
hello
```

Notice that sender got random, OS-assigned port

**Can a receiver written in C receive packets from a sender written in Python?
Yes!**

We also do not care if one computer is a Mac and the other is Windows

Promiscuous (monitor) mode sniffs all frames



Sniffing all frames can take a lot of CPU time!

sniff_raw.c

```
int main() {  
    int PACKET_LEN = 512;  
    char buffer[PACKET_LEN];  
    struct sockaddr saddr;  
    struct packet_mreq mr;  
  
    // Create the raw socket  
    int sock = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_ALL));  
  
    // Turn on the promiscuous mode.  
    mr.mr_type = PACKET_MR_PROMISC;  
    setsockopt(sock, SOL_PACKET, PACKET_ADD_MEMBERSHIP, &mr, sizeof(mr));  
  
    // Getting captured packets  
    while (1) {  
        int data_size = recvfrom(sock, buffer, PACKET_LEN, 0,  
                                &saddr, (socklen_t*)sizeof(saddr));  
        if(data_size) printf("Got one packet\n");  
    }  
  
    close(sock);  
    return 0;  
}
```

Sniff raw frames

Sniff all protocols (set to ETH_P_IP for only IP layer 3)

Turn on promiscuous mode to see all frames

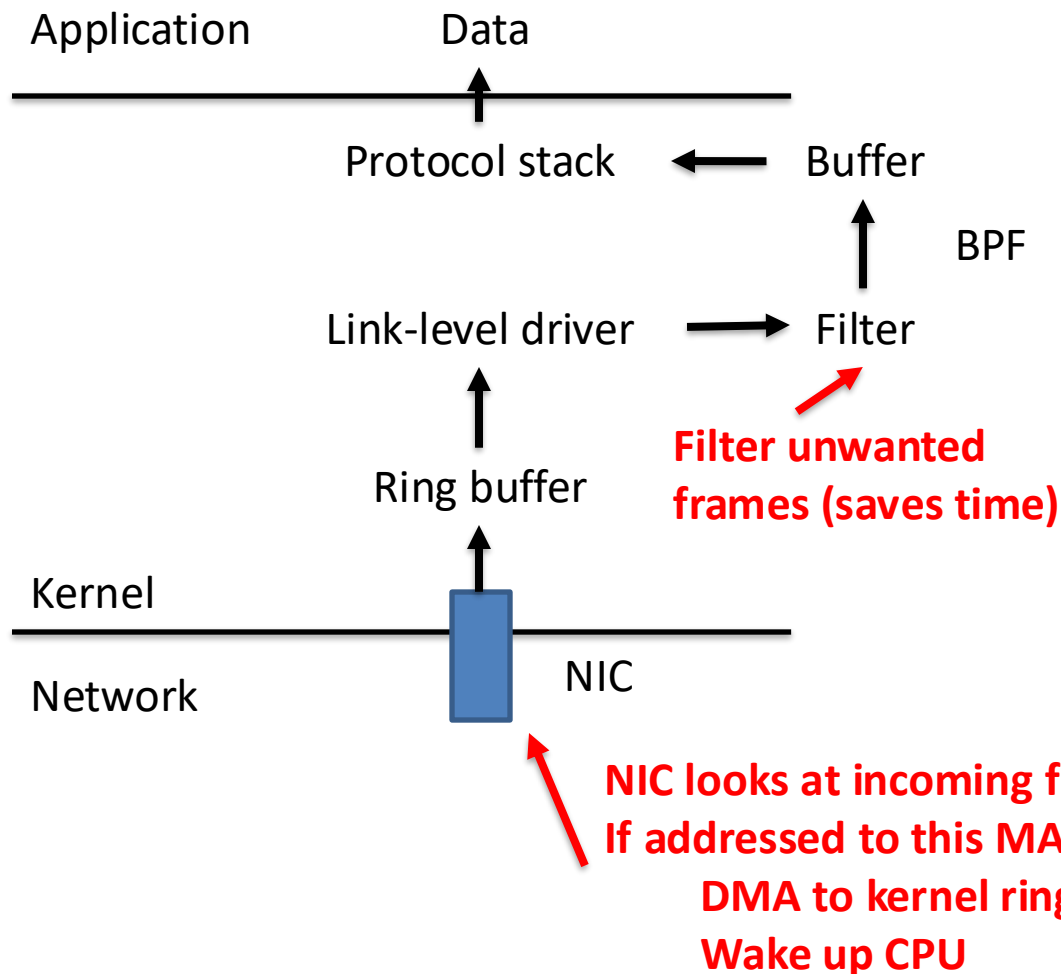
Set socket option enables promiscuous mode

Now see all layer 2 frames

Be careful, this may be very busy if you are on a high-traffic network!

Most of the time we will want to filter out unwanted frames to reduce processing

Layer 2 overview



Berkeley Packet Filter (BPF)

- Creates filter at link-level
- Only passes frames matching criteria to protocol stack
- Does not DMA to memory
- Compile filter and set with `setsockopt`

BPF is somewhat complicated
Add to raw socket with `SO_ATTACH_FILTER`
Not portable between OSes
PCAP API is much easier!

PCAP makes it easy to filter frames at a low level

PCAP (packet capture)

- Originally written for tcpdump (powerful sniffer)
- Supported by multiple platforms
 - Linux: libpcap
 - Windows: WinPcap and Npcap
- Written in C
- Other languages generally provide a wrapper around C version
- Basis used by other sniffers
 - Tcpdump (of course)
 - Wireshark
 - Scapy
 - Nmap
 - Snort

Use PCAP and compile BPF filter to filter out unwanted frames

```
void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet) {  
    printf("Got a packet\n");  
}
```

Prints got packet when UDP or ICMP packet arrives
Last parameter has packet details

```
int main() {  
    pcap_t *handle;  
    char errbuf[PCAP_ERRBUF_SIZE];  
    struct bpf_program fp;  
    char filter_exp[] = "udp or icmp";  
    bpf_u_int32 net;  
    const int MAX_SIZE = 8192;
```

Open PCAP session

- Sniff on interface *en0* (use ifconfig to find)
- 3rd parameter sets promiscuous mode to true

Parsing packets in C is tedious

Why?

Layer 2, 3, and 4 headers not constant size (optional components)

// Step 1: Open live pcap session on NIC with name enp0s3

```
handle = pcap_open_live("en0", MAX_SIZE, 1, 1000, errbuf);  
if (handle == NULL) { printf("Error on open\n"); printf("errbuf %s\n", errbuf); return (1); }
```

// Step 2: Compile filter_exp into BPF psuedo-code

```
pcap_compile(handle, &fp, filter_exp, 0, net);  
if (pcap_setfilter(handle, &fp) != 0) { printf("set filter error"); return(1); }
```

Compile BPF filter to pass
UDP or ICMP packets

// Step 3: Capture packets

```
pcap_loop(handle, -1, got_packet, NULL);
```

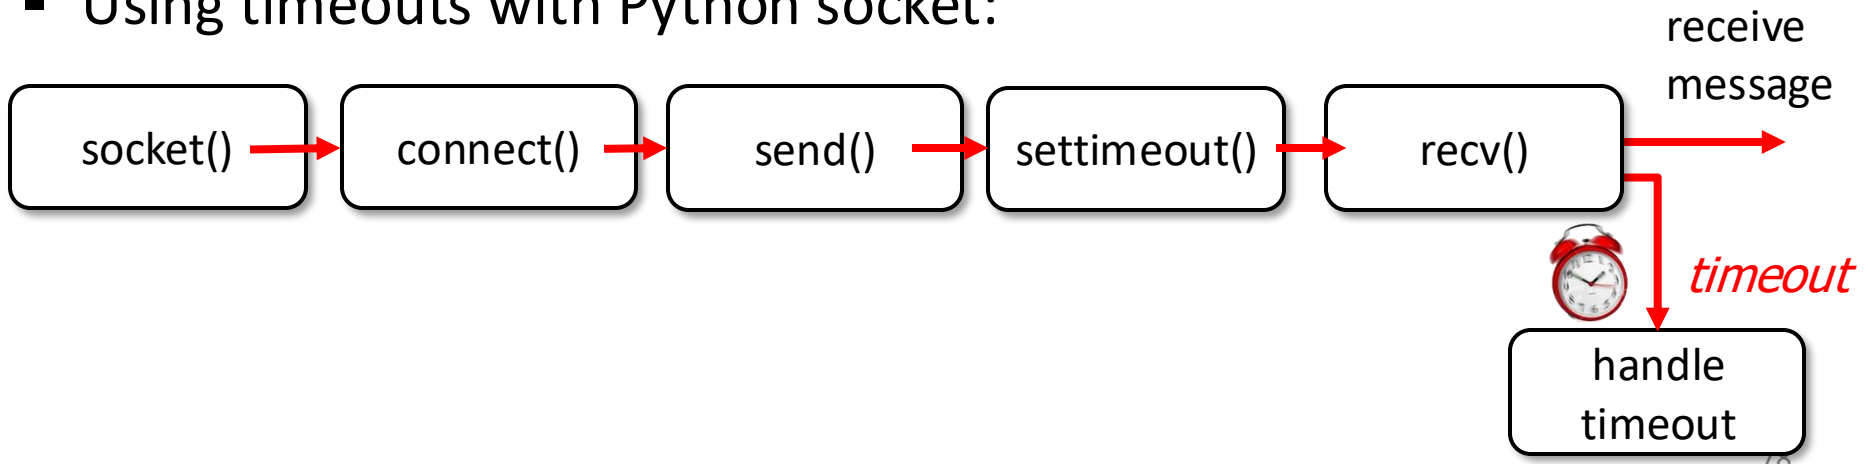
Start sniffing

Set callback function to *got_packet*

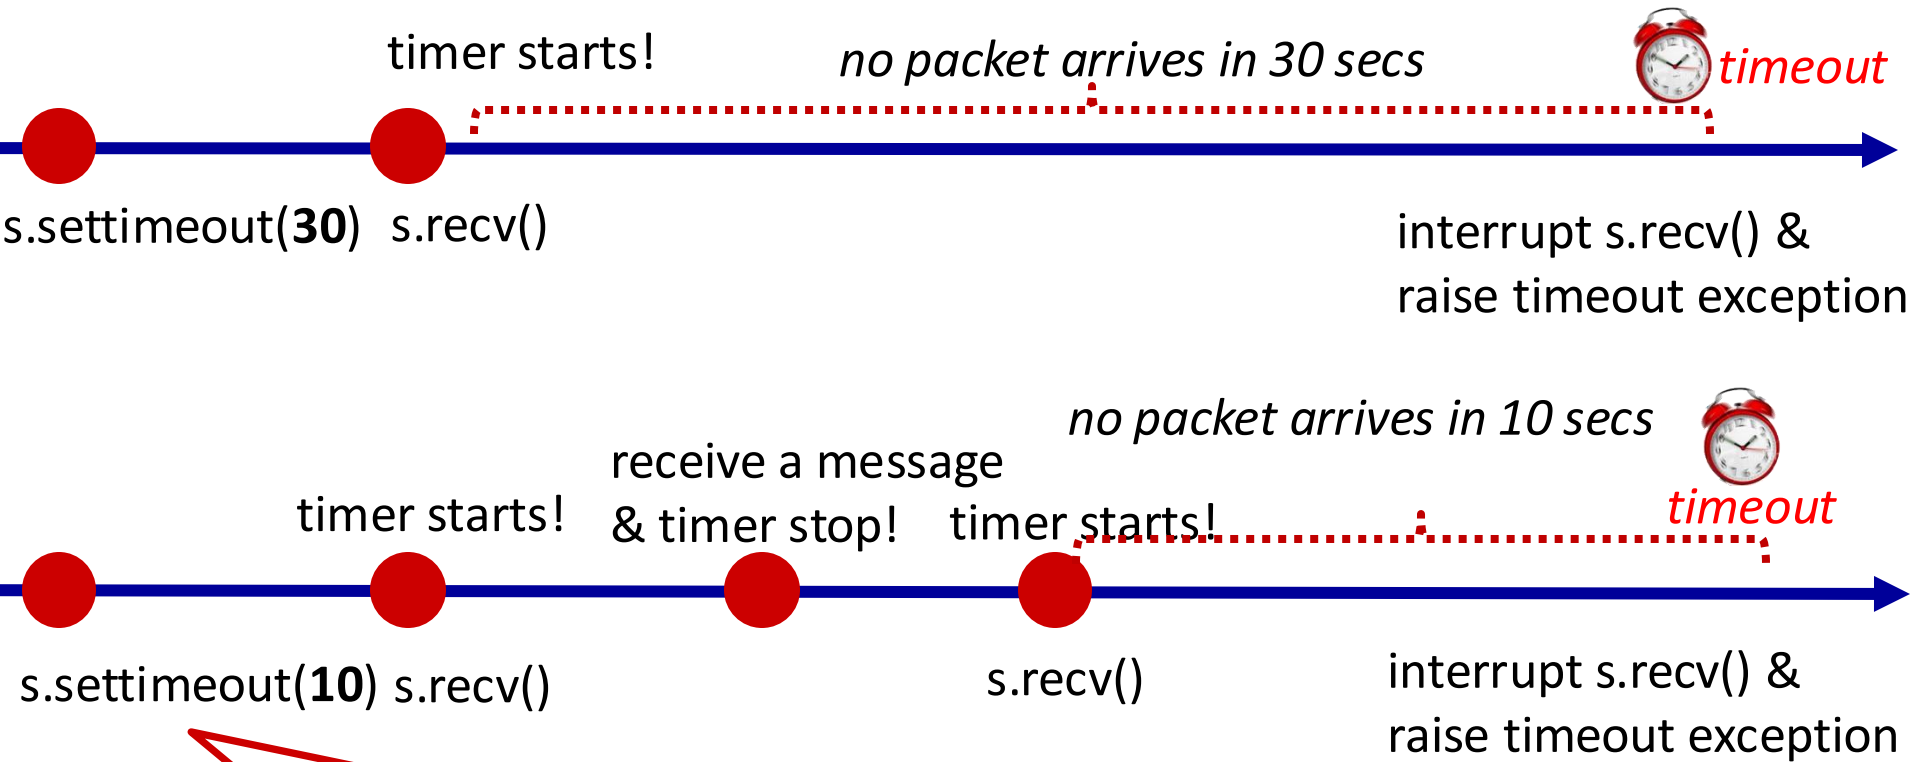
```
pcap_close(handle); //Close the handle  
return 0;
```

We can set timeouts on sockets

- Sometimes a program must **wait for one of several events** to happen, e.g.,:
 - Wait for either (i) a reply from another end of the socket, or (ii) timeout: **timer**
 - Wait for replies from several different open sockets: **select()**, **multithreading**
- Timeouts are used extensively in networking
- Using timeouts with Python socket:



Sockets timeout if no message is received



Set a timeout on all future socket operations of that specific socket!

Timeout exception raised if no response from server in 5 seconds

settimeout.py

```
# Create a UDP socket
```

```
sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
# Set a timeout of 5 seconds for the socket operations
```

```
sock.settimeout(5)
```

 **Set timeout to 5 seconds**

```
# Example usage (sending and receiving data with timeout)
```

```
serverAddressPort = ("127.0.0.1", 9090)
```

```
bytesToSend = str.encode("Hello UDP Server")
```

```
try:
```

```
    sock.sendto(bytesToSend, serverAddressPort)
```

```
    msgFromServer, address = sock.recvfrom(1024)
```

```
    print(f"Message from server: {msgFromServer.decode()}")
```

```
except socket.timeout:
```

```
    print("Socket operation timed out.")
```

```
except socket.error as e:
```

```
    print(f"Socket error: {e}")
```

```
finally:
```

```
    sock.close()
```

 **Begin timer on recvfrom
Throw exception if reply
not received in 5 seconds**

 **Catch exception thrown if no reply**

Agenda

1. Review: network layers

2. Sockets



3. Scapy

4. Exercises

Scapy makes it easy to sniff and transmit

Scapy

A Python-based packet manipulation program

Enables:

- Packet crafting (build any kind of packet)
- Sniffing (capture packets from the wire)
- Sending (inject custom packets)
- Dissection (decode fields of captured packets)

Works at Layers 2–7 (Ethernet to Application)

Takes care of details like calculating checksums for you!

Requires root privileges to run: `$ sudo python3 prog.py`

Install with: `$ pip3 install scapy`

Scapy uses Sockets under the hood but gives you control over header fields

Uses raw sockets for full control over network frames

Can list fields available at each network level with *ls(<layer name>)*

```
$ sudo python3
>>> from scapy.all import *
>>> ls(Ether)
dst : DestMACField = ('None')
src : SourceMACField = ('None')
type : XShortEnumField = ('36864')
>>> ls(IP)
version : BitField (4 bits) = ('4')
ihl : BitField (4 bits) = ('None')
tos : XByteField = ('0')
len : ShortField = ('None')
id : ShortField = ('1')
flags : FlagsField = ('<Flag 0 (>')
frag : BitField (13 bits) = ('0')
ttl : ByteField = ('64')
proto : ByteEnumField = ('0')
chksum : XShortField = ('None')
src : SourceIPField = ('None')
dst : DestIPField = ('None')
options : PacketListField = ('[]')
```

```
>>> ls(TCP)
sport : ShortEnumField = ('20')
dport : ShortEnumField = ('80')
seq : IntField = ('0')
ack : IntField = ('0')
dataofs : BitField (4 bits) = ('None')
reserved : BitField (3 bits) = ('0')
flags : FlagsField = ('<Flag 2 (S)>')
window : ShortField = ('8192')
chksum : XShortField = ('None')
urgptr : ShortField = ('0')
options : TCPOptionsField = ('b'')
>>>
```

Scapy calculates checksums and other fields for you!

Chooses “reasonable” values for fields

You can set or read any of these fields

Scapy makes it easy to use PCAP from Python to sniff packets with a filter

sniff_scapy.py

Scapy built on top of PCAP

```
from scapy.all import *
```

```
def process_packet(pkt):
```

← Called whenever packet arrives

```
    global count
```

```
    print('-'*40)
```

```
    print("Packet number",count)
```

```
    pkt.show()
```

← Print packet summary (much easier to process packets than with C)

```
    print('-'*40)
```

```
    count += 1
```

```
if __name__ == '__main__':
```

```
    count = 0
```

```
    pkt = sniff(iface='en0', filter='icmp or udp', count=10, prn=process_packet)
```

Sniff on this interface

Filter out packets other than these

Only capture 10 packets

Callback function when packet arrives

Run

```
$ sudo python3 sniff_scapy.py
```

← Uses promiscuous mode, so need sudo privilege

Another terminal

```
$ ping 8.8.8.8
```

Scapy makes it easy to use PCAP from Python to sniff packets with a filter

sniff_scapy.py

```
from scapy.all import *
```

```
def process_packet(pkt):
```

```
    global count
```

```
    print('-'*40)
```

```
    print("Packet number",count)
```

```
    pkt.show()
```

```
    print('-'*40)
```

```
    count += 1
```

```
if __name__ == '__main__':
```

```
    count = 0
```

```
    pkt = sniff(iface='en0', filter='icmp or udp')
```

Run

```
$ sudo python3 sniff_scapy.py
```

```
Packet number 0
###[ Ethernet ]###
  dst   = 00:50:56:f5:9e:21
  src   = 00:0c:29:a8:d0:96
  type  = IPv4
###[ IP ]###
  version = 4
  ihl     = 5
  tos     = 0x0
  len     = 84
  id      = 54952
  flags   = DF
  frag    = 0
  ttl     = 64
  proto   = icmp
  chksum  = 0xf3c6
  src     = 192.168.159.129
  dst     = 8.8.8.8
  \options \
###[ ICMP ]###
  type    = echo-request
  code    = 0
  chksum  = 0xa818
  id      = 0xb
  seq     = 0x1
  unused  = b''
###[ Raw ]###
  load    =
b'\x86\xb8\xd1h\x00\x00\x00\x00*\xe7\x0e\x00\x00\x00\x00\x00\x10\x
11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f
!"#$%&\'()*+,-./01234567'
```

Demo: packet spoofing is easy with Scapy

spoof_udp.py

False flag operation: it looks like this packet came from another computer!

We could have (but didn't here) create an Ether layer and spoof Layer 2 MAC addresses

```
#!/usr/bin/python3
```

```
from scapy.all import *
```

You will do this in Lab 1

```
print("SENDING SPOOFED UDP PACKET.....")
```

```
ip = IP(src="1.2.3.4", dst="10.0.2.5") # IP Layer
```

Create IP layer 3 with spoofed src and dest

```
udp = UDP(sport=8888, dport=9090) # UDP Layer
```

```
data = "Hello UDP!\n" # Payload
```

```
pkt = ip/udp/data # Construct the complete packet
```

```
pkt.show()
```

```
send(pkt, verbose=0)
```

Set UDP (layer 4) source and destination ports to destination

Send packet

Stack layers to create UDP over IP layer with data

Scapy fills in default values and calculates checksums for us!
Nice!

We could do this in C also, but tedious to set packet values

Demo: packet spoofing is easy with Scapy

spoof_udp.py

Terminal 1 (sender)

```
$ sudo python3 spoof_udp.py
```

```
#!/usr/bin/python3
```

```
from scapy.all import *
```

```
print("SENDING SPOOFED UDP PACKET.....")
```

```
ip = IP(src="1.2.3.4", dst="10.0.2.5") # IP Layer
```

```
udp = UDP(sport=8888, dport=9090) # UDP Layer
```

```
data = "Hello UDP!\n" # Payload
```

```
pkt = ip/udp/data # Construct the complete packet
```

```
pkt.show()
```

```
send(pkt, verbose=0)
```

Terminal 2 (sender)

```
$ sudo python3 sniff_scapy.py ens160
```

```
####[ Ethernet ]####
```

```
dst      = 00:50:56:f5:9e:21
```

```
src      = 00:0c:29:a8:d0:96
```

```
type     = IPv4
```

```
####[ IP ]####
```

```
version  = 4
```

```
ihl      = 5
```

```
tos      = 0x0
```

```
len      = 39
```

```
id       = 1
```

```
flags    =
```

```
frag     = 0
```

```
ttl      = 64
```

```
proto    = udp
```

```
chksum   = 0x6abb
```

```
src      = 1.2.3.4
```

```
dst      = 10.0.2.5
```

```
\options \
```

```
####[ UDP ]####
```

```
sport    = 8888
```

```
dport    = 9090
```

```
len      = 19
```

```
chksum   = 0xd62b
```

```
####[ Raw ]####
```

```
load     = b'Hello UDP!\n'
```

Sometimes we want to sniff, then spoof a reply

sniff_spoof_icmp.py

```
from scapy.all import *
```

```
def spoof_pkt(pkt):
```

```
    if ICMP in pkt and pkt[ICMP].type == 8:  
        #listen for ICMP request packets (type 8)
```

```
        print("Original Packet.....")
```

```
        print("Source IP : ", pkt[IP].src)
```

```
        print("Destination IP :", pkt[IP].dst)
```

```
        #spoof a reply, even if the request wasn't for us  
        #must reverse source and destination on reply!
```

```
        ip = IP(src=pkt[IP].dst, dst=pkt[IP].src, ihl=pkt[IP].ihl)
```

```
        icmp = ICMP(type=0, id=pkt[ICMP].id, seq=pkt[ICMP].seq)
```

```
        data = pkt[Raw].load
```

```
        newpkt = ip/icmp/data
```

```
        print("Spoofed Packet.....")
```

```
        print("Source IP : ", newpkt[IP].src)
```

```
        print("Destination IP :", newpkt[IP].dst)
```

```
        send(newpkt, verbose=0)
```

Create a phantom computer on the network that responds to pings from 10.0.2.15!

Steps

- Receive packet
- Extract details
- Spoof a new packet
 - Reverse src and dst
 - Keep same id and seq number

Send a ICMP reply after ping request to 10.0.2.15

```
pkt = sniff(filter='icmp and src host 10.0.2.15', prn=spoof_pkt)
```


Scapy vs. C

Scapy

Pros

- Easier to write
- Sets reasonable default fields
- Calculates values (checksums)
- Focus on fields interested
- Every layer is an object, can easily stack together
- Increased productivity

C

Pros

- Runs much faster! (50-100X!)
- More control over crafting packets

Scapy vs. C

Scapy

Pros

- Easier to write
- Sets reasonable default fields
- Calculates values (checksums)
- Focus on fields interested
- Every layer is an object, can easily stack together
- Increased productivity

Cons

- Runs slowly

C

Pros

- Runs much faster! (50-100X!)
- More control over crafting packets

Cons

- Tricky to write, must get the byte offsets right
- Tedious?

Scapy and C can work together in a hybrid approach

Hybrid approach

For some uses speed is critical

Scapy is slow

Example:

Send 1000 UDP packets

- Scapy: 9.4 seconds
- C : 0.25 seconds
- C is 37X faster than Scapy in this case!



Sometimes racing to reply before the “real” computer – Scapy too slow

Idea:

- Use Scapy to create packets, and save to file
- Read packets in C (might to adjust some fields)
- Send using C

Agenda

1. Review: network layers

2. Sockets

3. Scapy

 4. Exercises

Exercise

Implement a chat program where:

- VPN to Dartmouth's network first (Dartmouth blocks traffic)!
- The client takes input from the keyboard and sends it to the server over UDP
- Server responds back to client with message converted to all uppercase
- Start with `send_udp.py/receive_udp.py`

Time permitting

- Implement the same in C
- Start with `send_udp.c/receive_udp.c`

Endianness

Endianness: a term that refers to the order in which a given multi-byte data item is stored in memory

- **Little Endian**: store the most significant byte of data at the highest address
- **Big Endian**: store the most significant byte of data at the lowest address

**Little endian
always seems
backwards to me**

Store 0xDEADBEEF

Big endian



Increasing addresses →

Little endian



Endianness in network communication

Computers with different byte orders will misunderstand each other

- Solution: agree upon a common order for communication
- This is called “network order”, which is the same as Big Endian order
- But Intel computers (“hosts”) use Little Endian
- Must convert data between “host order” and “network order”

Macro	Description
<code>htons()</code>	Convert unsigned short integer from host order to network order.
<code>htonl()</code>	Convert unsigned integer from host order to network order.
<code>ntohs()</code>	Convert unsigned short integer from network order to host order.
<code>ntohl()</code>	Convert unsigned integer from network order to host order.

Prior to 1993 IP addresses were assigned in address range blocks

Classful addressing scheme (1981 -1993)

Class A: 0.0.0.0 – 127.255.255.255

Used to help routing,
32 bits total

Left bit always 0, if see you a
zero in the left bit, you know it
belongs to a Class A network

Class B: 128.0.0.0 – 191.255.255.255

Class C: 192.0.0.0 – 223.255.255.255

Network ID is $2^7 = 127$ possible
address for first octet

Class D: 224.0.0.0 – 239.255.255.255

If you had Class A network
(first 8 bits set) you have $2^{24} =$
16.7M possible addresses that
can be given to devices

Class E: 240.0.0.0 – 255.255.255.255

Issued to very large
organizations

Prior to 1993 IP addresses were assigned in address range blocks

Classful addressing scheme (1981 -1993)

Class A: 0.0.0.0 – 127.255.255.255

Class B: 128.0.0.0 – 191.255.255.255

Class C: 192.0.0.0 – 223.255.255.255

Class D: 224.0.0.0 – 239.255.255.255

Class E: 240.0.0.0 – 255.255.255.255

Used to help routing
32 bits total

Network ID is left 16 bits

- Left bit = 1
- Left second bit = 0

If you see left bits are 10, you know this is a Class B network

If you had Class B network (first 16 bits set) you have $2^{16} = 65.5K$ possible addresses that can be given to devices

Issued to large organizations

Prior to 1993 IP addresses were assigned in address range blocks

Classful addressing scheme (1981 -1993)

Class A: 0.0.0.0 – 127.255.255.255

Class B: 128.0.0.0 – 191.255.255.255

Class C: 192.0.0.0 – 223.255.255.255

Class D: 224.0.0.0 – 239.255.255.255

Class E: 240.0.0.0 – 255.255.255.255

Used to help routing
32 bits total

Network ID is left 24 bits

- Left bit = 1
- Left second bit = 1
- Third left bit = 0

If you see left bits are 110, you know this is a Class C network

If you had Class C network (first 24 bits set) you have $2^8 = 256$ possible addresses that can be given to devices

Issued to medium sized organizations

Prior to 1993 IP addresses were assigned in address range blocks

Classful addressing scheme (1981 -1993)

Class A: 0.0.0.0 – 127.255.255.255

Class B: 128.0.0.0 – 191.255.255.255

Class C: 192.0.0.0 – 223.255.255.255

Class D: 224.0.0.0 – 239.255.255.255

Class E: 240.0.0.0 – 255.255.255.255

Used to help routing
32 bits total

Class D for multicast (special purpose)

Class E reserved

Problem:

We were using up IP addresses too quickly (lots of unused addresses if you own a Class A network)

Happened to us, we had a class C and had to give up some addresses

